

Processor Architecture

Victor Eijkhout

Fall 2025

Justification

The performance of a parallel code has as one component the behaviour of the single processor or single-threaded code. In this section we discuss the basics of how a processor executes instructions, and how it handles the data these instructions operate on.

Structure of a modern processor

Von Neumann machine

The ideal processor:

- ▶ (Stored program)
- ▶ An instruction contains the operation and two operand locations
- ▶ Processor decodes instruction, gets operands, computes and writes back the result
- ▶ Repeat

The actual state of affairs

- ▶ Single instruction stream versus multiple cores / floating point units
- ▶ Single instruction stream versus Instruction Level Parallelism
- ▶ Unit-time-addressable memory versus large latencies

Modern processors contain lots of magic to make them seem like Von Neumann machines.

Complexity measures

Traditional: processor speed was paramount. Operation counting.


Nowadays: memory is slower than processors

This lecture:

Study data movement aspects

Algorithm design for processor reality

A first look at a processor



Copyright (c) 2011 Intel Corporation. All rights reserved.

Structure of a core

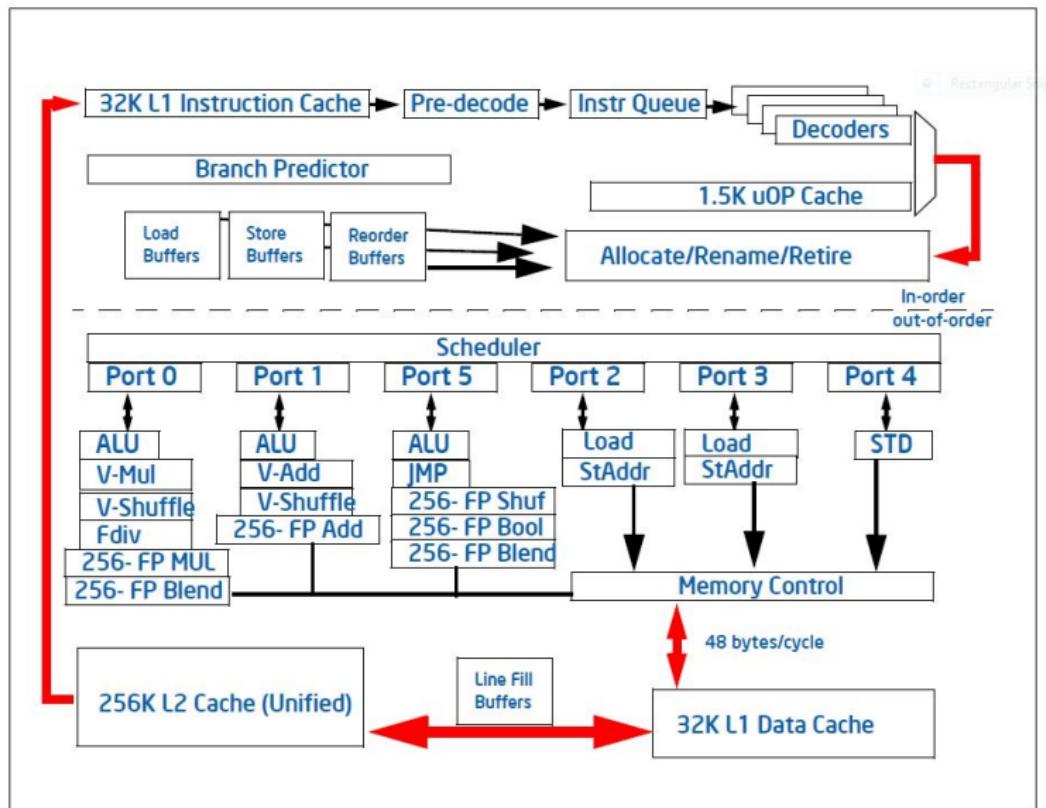


Figure 2-1. Intel microarchitecture code name Sandy Bridge Pipeline Functionality

Motivation for pipelining

An operation consists of several stages.

Addition:

- ▶ Decoding the instruction operands.
- ▶ Data fetch into register
- ▶ Aligning the exponents:

$$\begin{aligned} .35 \times 10^{-1} + .6 \times 10^{-2} &\quad \text{becomes} \\ .35 \times 10^{-1} + .06 \times 10^{-1}. \end{aligned}$$





- ▶ Adding mantissas, giving .41.
- ▶ Normalizing the result, giving $.41 \times 10^{-1}$.
- ▶ Storing the result.

pipeline stages

Pipelining, pictorially

Discrete hardware for each stage:

$$c_i \leftarrow a_i + b_i$$



Analysis

Operation timing:

$$\begin{cases} n & \text{operations} \\ \ell & \text{number of stages} \Rightarrow t(n) = n\ell\tau \\ \tau & \text{clock cycle} \end{cases}$$

With pipelining:

$$t(n) = [s + \ell + n - 1]\tau$$

where s is a setup cost

\Rightarrow Asymptotic speedup is ℓ

$n_{1/2}$: value for which speedup is $\ell/2$

Applicability of pipelining

Pipelining works for:

vector addition/multiplication

Division/square root maybe pipelined, but much slower

Recurrences

Pipelining does not immediately work:

```
for (i) {  
    x[i+1] = a[i]*x[i] + b[i];  
}
```

Transform:

$$\begin{aligned}x_{n+2} &= a_{n+1}x_{n+1} + b_{n+1} \\&= a_{n+1}(a_nx_n + b_n) + b_{n+1} \\&= a_{n+1}a_nx_n + a_{n+1}b_n + b_{n+1}\end{aligned}$$

Instruction pipeline

- ▶ Instruction-Level Parallelism: more general notion of independent instructions
- ▶ Requires independent instructions
- ▶ As frequency goes up, pipeline gets longer: more demands on compiler

Instruction-Level Parallelism

- ▶ multiple-issue of independent instructions
- ▶ branch prediction and speculative execution
- ▶ out-of-order execution
- ▶ prefetching

Problems: complicated circuitry, hard to maintain performance

Implications

- ▶ Long pipeline needs many independent instructions:
demands on compiler
- ▶ Conditionals break the stream of independent instructions
 - ▶ Processor tries to predict branches
 - ▶ branch misprediction penalty:
pipeline needs to be flushed and refilled
 - ▶ avoid conditionals in inner loops!

Instructions

- ▶ Addition/multiplication: pipelined
- ▶ Division (and square root): much slower

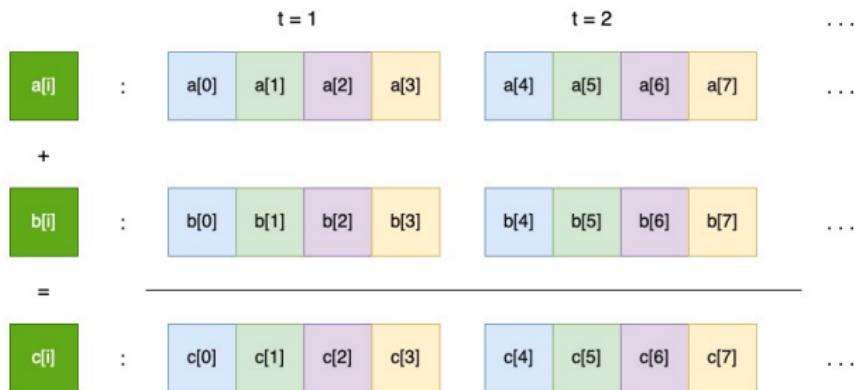
```
1 for ( i )  
2     a[i] = b[i] / c
```

Can you improve on this?

- ▶ Fused Multiply-Add (FMA) $s += a*b$
where can you use this?

Vector instructions

- ▶ Multiple (independent) operands per cycle
 - ▶ Similar to pipelining but not the same
 - ▶ ‘Vector width’ is 2,4,8.



DAXPY

BLAS operation ‘A-x plus y’:

```
1 for ( i )  
2     y[i] = a[i]*x + y[i]
```

- ▶ x stays in register
- ▶ two loads and one store per operation
- ▶ candidate for pipelining or vector instructions

DAXPY performance

Performance is a function of

- ▶ Clock frequency,
- ▶ SIMD width
- ▶ Load/store unit behavior

Floating point capabilities of several processor architectures

DAXPY cycle number for 8 operands

Processor	year	add/mult/fma units (count \times width)	daxpy cycles (arith vs load/store)
MIPS R10000	1996	$1 \times 1 + 1 \times 1 + 0$	8/24
Alpha EV5	1996	$1 \times 1 + 1 \times 1 + 0$	8/12
IBM Power5	2004	$0 + 0 + 2 \times 1$	4/12
AMD Bulldozer	2011	$2 \times 2 + 2 \times 2 + 0$	2/4
Intel Sandy Bridge	2012	$1 \times 4 + 1 \times 4 + 0$	2/4
Intel Haswell	2014	$0 + 0 + 2 \times 4$	1/2

Memory hierarchy: caches, register, TLB.

The Big Story


- ▶ DRAM memory is slow, so let's put small SRAM close to the processor
- ▶ This helps if data is reused: second access is faster
- ▶ Does the algorithm have reuse?
- ▶ Does the implementation reuse data?

Bandwidth and latency

Important theoretical concept:

- ▶ latency is delay between request for data and availability
- ▶ bandwidth is rate at which data arrives thereafter

Memory hierarchy



Registers

Computing out of registers

a := b + c

- ▶ load the value of b from memory into a register,
- ▶ load the value of c from memory into another register,
- ▶ compute the sum and write that into yet another register, and
- ▶ write the sum value back to the memory location of a.

Register usage

Assembly code

(note: Intel two-operand syntax)

```
addl %eax, %edx
```

- ▶ Registers are named
- ▶ Can be explicitly addressed by the programmer
- ▶ ... as opposed to caches.
- ▶ Assembly coding or inline assembly (compiler dependent)
- ▶ ... but typically generated by compiler

Examples of register usage

1. Resident in register

```
a := b + c
```

```
d := a + e
```

a stays resident in register, avoid store and load

2. subexpression elimination:

```
t1 = sin(alpha) * x + cos(alpha) * y;
```

```
t2 = -cos(alpha) * x + sin(alpha) * y;
```

becomes:

```
s = sin(alpha); c = cos(alpha);
```

```
t1 = s * x + c * y;
```

```
t2 = -c * x + s * y
```

often done by compiler

Caches

Cache basics

Fast SRAM in between memory and registers: mostly serves data reuse

```
... = ... x ..... // instruction using x  
.....           // several instructions not involving x  
... = ... x ..... // instruction using x
```

- ▶ load x from memory into cache, and from cache into register; operate on it;
- ▶ do the intervening instructions;
- ▶ request x from memory, but since it is still in the cache, load it from the cache into register; operate on it.
- ▶ essential concept: data reuse

Cache levels

- ▶ Levels 1,2,3(,4): L1, L2, etc.
- ▶ Increasing size, increasing latency, decreasing bandwidth
- ▶ (Note: L3/L4 can be fairly big; beware benchmarking)
- ▶ Cache hit / cache miss: one level is consulted, then the next
- ▶ L1 has separate data / instruction cache, other levels mixed
- ▶ Caches do not have enough bandwidth to serve the processor:
coding for reuse on all levels.

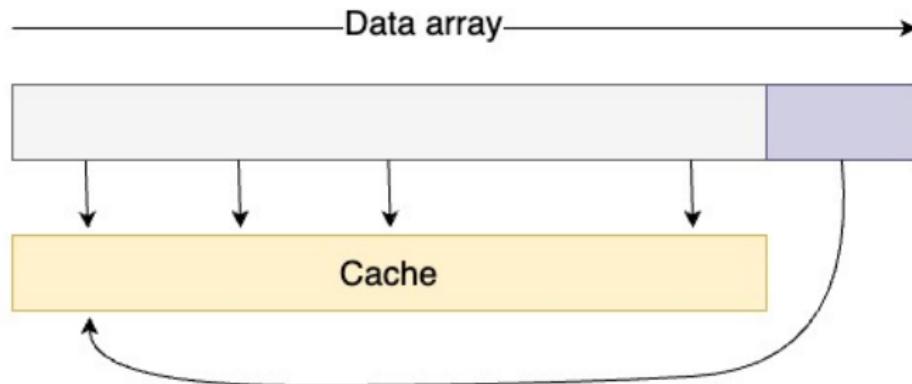
Cache misses

- ▶ Compulsory miss: first time data is referenced
- ▶ Capacity miss: data was in cache, but has been flushed (overwritten) by LRU policy
- ▶ Conflict miss: two items get mapped to the same cache location, even if there are no capacity problems
- ▶ Invalidation miss: data becomes invalid because of activity of another core

Cache hits

- ▶ Data has been requested, used a second time: temporal locality
- ▶ ⇒ Can't wait too long between uses
- ▶ (Data can be loaded because it's close to data requested: spatial locality. Later.)

Capacity miss



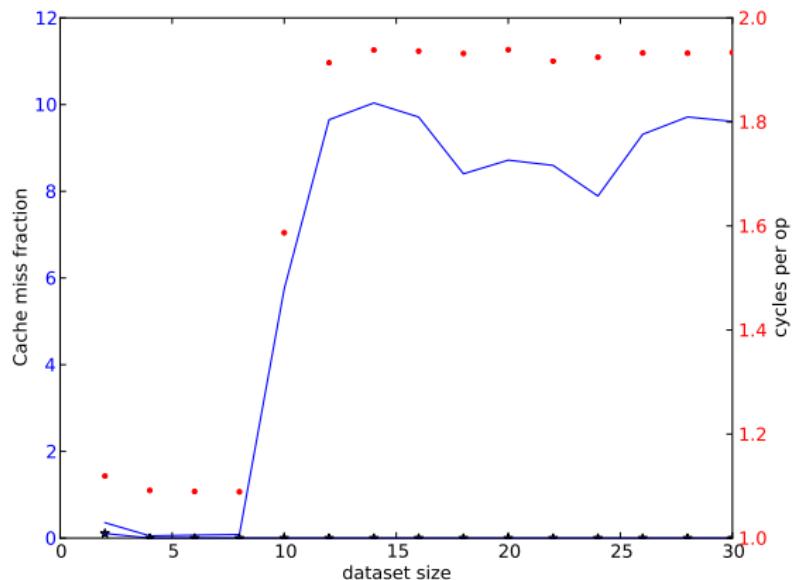
(Why is that last block going where it is going?)

Cache capacity

- ▶ Loading data multiple times
- ▶ LRU: oldest item evicted if needed
- ▶ Reuse if not too much data

```
1 for ( lots of times ) // sequential loop
2     load and process data // probably parallel loop
```

Illustration of capacity



Replacement policies

What determines where new data goes /
what old data is overwritten?

- ▶ Least Recently Used (LRU): most common
- ▶ First-in-first-out (FIFO): IBM Power4. Not a good idea.
- ▶ Random Replacement. Sometimes used.

It's actually more subtle than pure LRU . . .

Cache lines

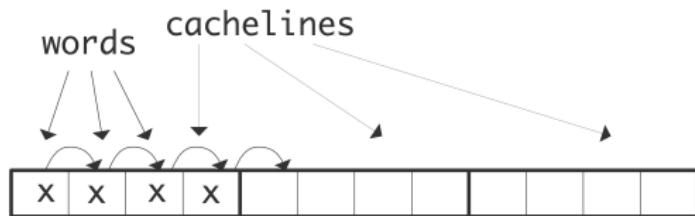
- ▶ Memory requests go by byte or word
- ▶ Memory transfers go by cache line:
typically 64 bytes / 8 double precision numbers
- ▶ Cache line transfer costs bandwidth
- ▶ ⇒ important to use all elements

Effects of striding

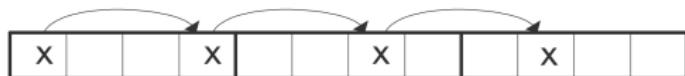
- ▶ Always 8 numbers transferred
- ▶ With stride $s > 1$: $8/s$ elements used
- ▶ Loss of efficiency if bandwidth-limited

Cache line use

```
for (i=0; i<N; i++)  
... = ... x[i] ...
```

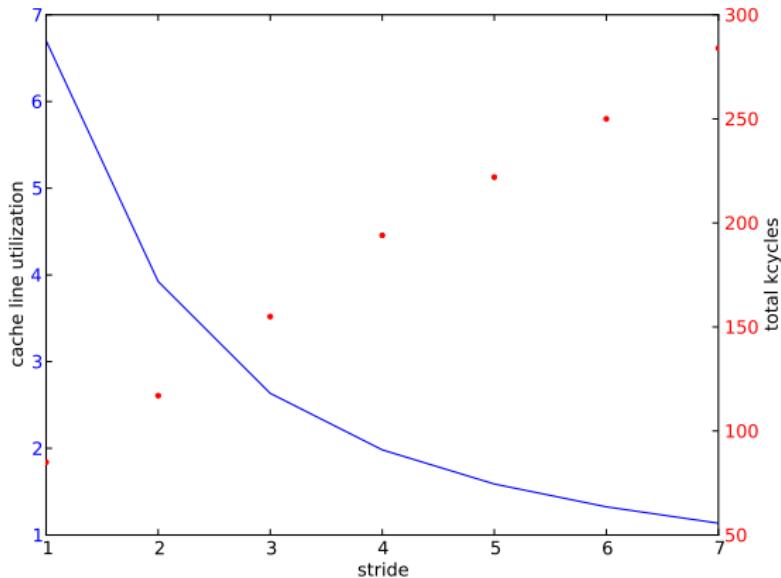


```
for (i=0; i<N; i+=stride)  
... = ... x[i] ...
```



Stride effects

```
for (i=0, n=0; i<L1WORDS; i++, n+=stride)
    array[n] = 2.3*array[n]+1.2;
```

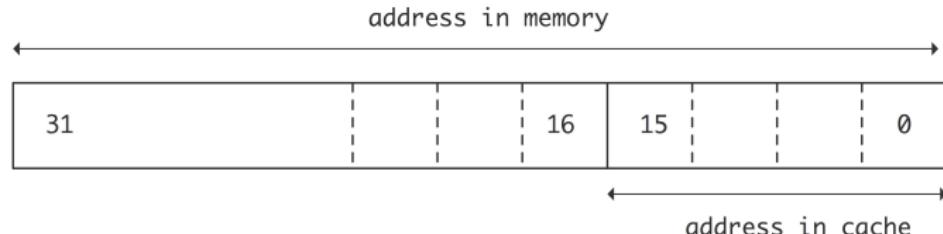


Cache mapping

Cache is smaller than memory, so we need a mapping scheme
memory address \mapsto cache address

- ▶ Ideal: any address can go anywhere; LRU policy for replacement
- ▶ pro: optimal; con: slow, expensive to manufacture
- ▶ Simple: direct mapping by truncating addresses
- ▶ pro: fast and cheap; con: I'll show you in a minute
- ▶ Practical: limited associativity
'enough but not too much'

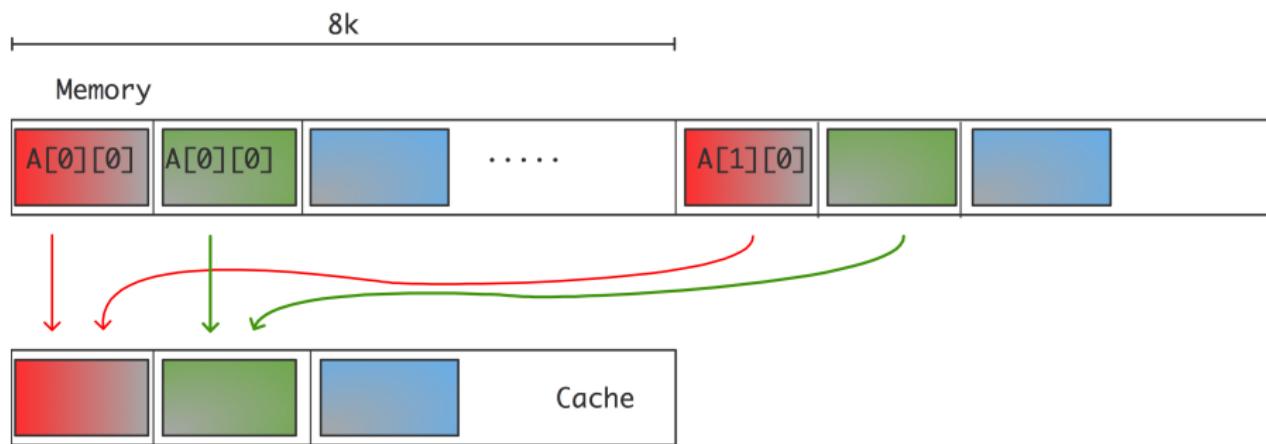
Direct mapping



Direct mapping of 32-bit addresses into a 64K cache

- ▶ Use last number of bits to find cache address
- ▶ If you traverse an array, a contiguous chunk will be mapped to cache without conflict.
- ▶ If (memory) addresses are cache size apart, they get mapped to the same cache location

Conflicts



Mapping conflicts in a direct-mapped cache.

The problem with direct mapping

```
real*8 A(8192,3);
do i=1,512
  a(i,3) = ( a(i,1)+a(i,2) )/2
end do
```

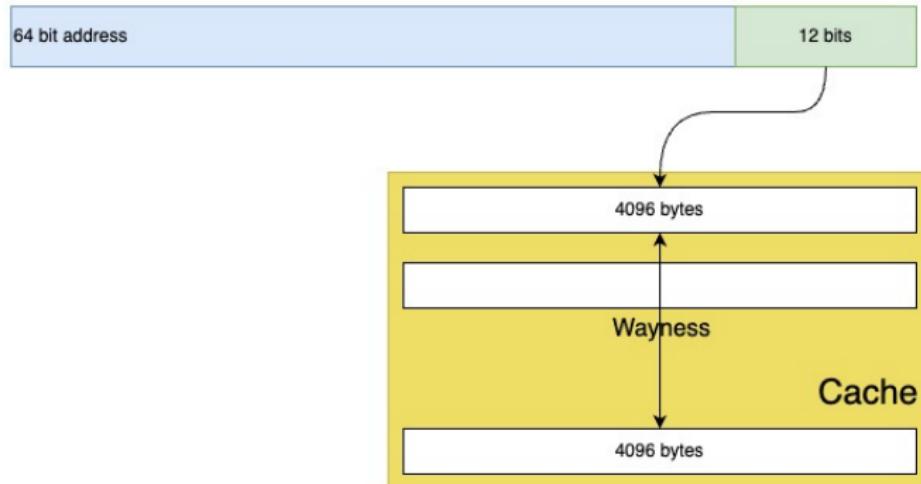
In each iteration 3 elements map to the same cache location:
constant overwriting ('eviction', cache thrasing):
low performance

Associative cache mapping

- ▶ Allow each memory address to go to multiple (but not all) cache addresses; typically 2,4,8
- ▶ Prevents problems with multiple arrays
- ▶ Reasonable fast
- ▶ Often lower associativity for L1 than L2, L3

Associativity	L1	L2
Intel (Woodcrest)	8	8
AMD (Bulldozer)	2	8

Associativity



Associative cache structure of Intel Ice Lake.

Illustration of associativity

{0, 12, 24, ... }
{1, 13, 25, ... }
{2, 14, 26, ... }
{3, 15, 27, ... }
{4, 16, 28, ... }
{5, 17, 29, ... }
{6, 18, 30, ... }
{7, 19, 31, ... }
{8, 20, 32, ... }
{9, 21, 33, ... }
{10, 22, 34, ... }
{11, 23, 35, ... }

{0, 12, 24, ... } {4, 16, 28, ... }
 {8, 20, 32, ... }
{1, 13, 25, ... } {5, 17, 29, ... }
 {9, 21, 33, ... }
{2, 14, 26, ... } {6, 18, 30, ... }
 {10, 22, 34, ... }
{3, 15, 27, ... } {7, 19, 31, ... }
 {11, 23, 35, ... }


Two caches of 12 elements: direct mapped (left) and 3-way associative (right)

Direct map: 0–12 is conflict

Associative: no conflict

Associativity in practice


$$\forall_j: y_j = y_j + \sum_{i=1}^m x_{i,j}$$



The number of L1 cache misses and the number of cycles for each j column accumulation, vector length 4096

One remedy

Do not user powers of 2.



The number of L1 cache misses and the number of cycles for each j column accumulation, vector length $4096 + 8$

Exercise

Write a small cache simulator in your favorite language. Assume a k -way associative cache of 32 entries and an architecture with 16 bit addresses.

Exercise: vectorsum

- ▶ Compare sequential performance to single-threaded OMP
- ▶ For some problem sizes observe a difference in performance
- ▶ Use Intel option `-qopt-report=3` and inspect the report.
- ▶ Compare different compilers: Intel 19 behaves differently from 24!
Also gcc13.

```
1 for ( int iloop=0; iloop<nloops; ++iloop ) {  
2     for ( int i=0; i<vectorsize; ++i ) {  
3         outvec[i] += invec[i]*loopcoeff[iloop];  
4     }  
5 }
```

Analyze and report

More memory system topics

Bandwidth / latency

Simple model for sending n words:

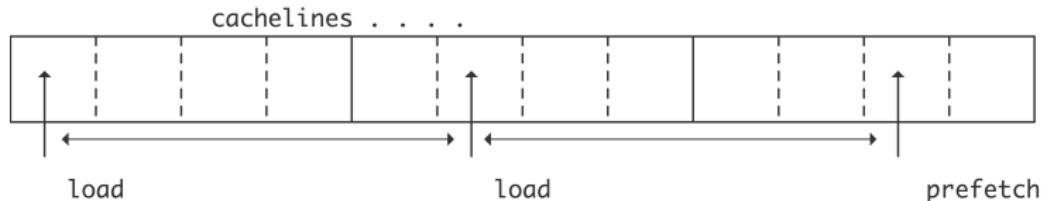
$$t = \alpha + \beta n$$

Quoted bandwidth figures are always optimistic:

- ▶ bandwidth shared between cores
not enough bandwidth for all cores:
 \Rightarrow speedup less than linear
- ▶ bandwidth wasted on coherence
- ▶ NUMA: pulling data from other socket
- ▶ assumes optimal scheduling of DRAM banks

Prefetch

- ▶ Do you have to wait for every item from memory?
- ▶ Memory controller can infer streams: prefetch
- ▶ Sometimes controllable through assembly, directives, libraries (AltiVec)
- ▶ One form of latency hiding



Memory pages

Memory is organized in pages:


- ▶ Translation between logical address, as used by program, and physical in memory
- ▶ This serves virtual memory and relocatable code
- ▶ so we need another translation stage.

Page translation: TLB

- ▶ General page translation: slowish and expensive
- ▶ Translation Look-aside Buffer (TLB) is a small list of frequently used pages
- ▶ Example of spatial locality: items on an already referenced page are found faster

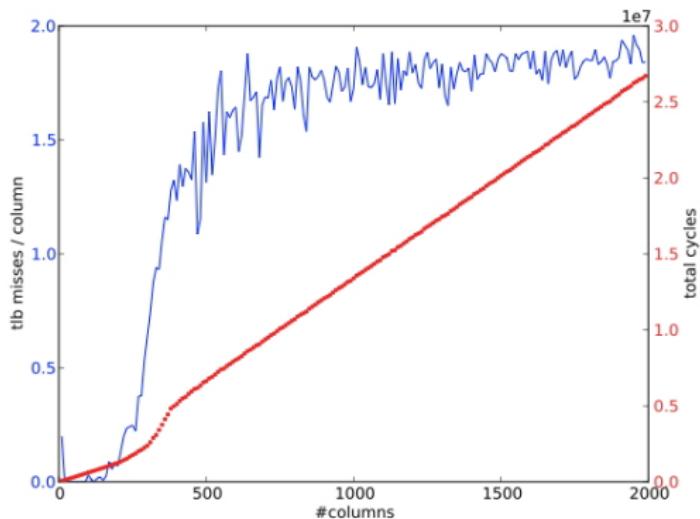
TLB misses

```
#define INDEX(i,j,m,n) i+j*m
array = (double*) malloc(m*n*sizeof(double));
/* traversal #2 */
for (i=0; i<m; i++)
    for (j=0; j<n; j++)
        array[INDEX(i,j,m,n)] = array[INDEX(i,j,m,n)]+1;
```



TLB hits


```
#define INDEX(i,j,m,n) i+j*m
array = (double*) malloc(m*n*sizeof(double));
/* traversal #1 */
for (j=0; j<n; j++)
    for (i=0; i<m; i++)
        array[INDEX(i,j,m,n)] = array[INDEX(i,j,m,n)]+1;
```



Little's Law

- ▶ Item loaded from memory, processed, new item loaded in response
- ▶ But this can only happen after latency wait
- ▶ Items during latency are independent, therefore

$$\text{Concurrency} = \text{Bandwidth} \times \text{Latency}.$$



Multicore issues

Why multicore

Quest for higher performance:

- ▶ Not enough instruction parallelism for long pipelines
- ▶ Two cores at half speed more energy-efficient than one at full speed.

Multicore solution:

- ▶ More theoretical performance
- ▶ Burden for parallelism is now on the programmer

Dennard scaling

Scale down feature size by s :

Feature size	$\sim s$
Voltage	$\sim s$
Current	$\sim s$
Frequency	$\sim s^{-1}$

Miracle conclusion:

$$\text{Power} = V \cdot I \sim s^2; \text{Power density} \sim 1$$

Everything gets better, cooling problem stays the same
Opportunity for more components, higher frequency

Dynamic power

Charge	$q = CV$
Work	$W = qV = CV^2$
Power	$W/\text{time} = WF = CV^2F$


(1)

Two cores at half frequency:

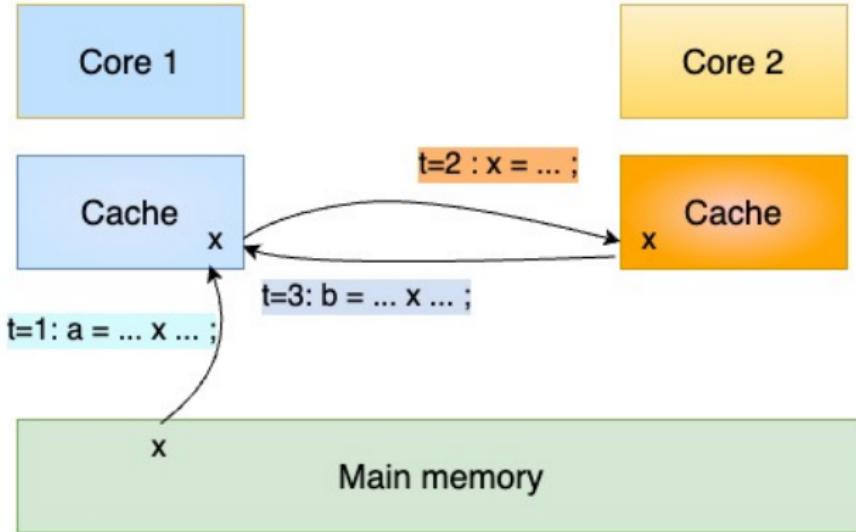
$$\left. \begin{array}{l} C_{\text{multi}} = 2C \\ F_{\text{multi}} = F/2 \\ V_{\text{multi}} = V/2 \end{array} \right\} \Rightarrow P_{\text{multi}} = P/4.$$

Same computation, less power

Multicore caches



The coherence problem



Cache coherence

Modified-Shared-Invalid (MSI) coherence protocol:

Modified: the cacheline has been modified

Shared: the line is present in at least one cache and is unmodified.

Invalid: the line is not present, or it is present but a copy in another cache has been modified.

Coherence issues

- ▶ Coherence is automatic, so you don't have to worry about it...
- ▶ ... except when it saps performance
- ▶ Beware false sharing
 - writes to different elements of a cache line

Balance analysis

- ▶ Sandy Bridge core can absorb 300 GB/s
- ▶ 4 DDR3/1600 channels provide 51 GB/s, difference has to come from reuse
- ▶ It gets worse: latency 80ns, bandwidth 51 GB/s,
Little's law: parallelism 64 cache lines
- ▶ However, each core only has 10 line fill buffers,
so we need 6–7 cores to provide the data for one core
- ▶ Power: cores are 72%, uncore 17, DRAM 11.
- ▶ Core power goes 40% to instruction handling, not arithmetic
- ▶ Time for a redesign of processors and programming; see my research presentation

Programming strategies for performance

How much performance is possible?

Performance limited by

- ▶ Processor peak performance: absolute limit
- ▶ Bandwidth: linear correlation with performance

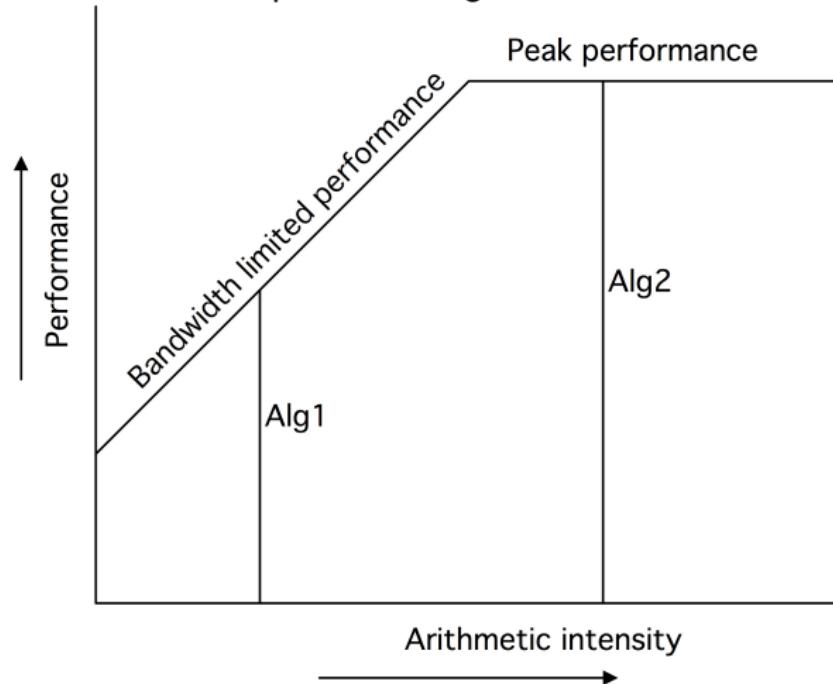
Arithmetic intensity: ratio of operations per transfer

If AI high enough: processor-limited

otherwise: bandwidth-limited

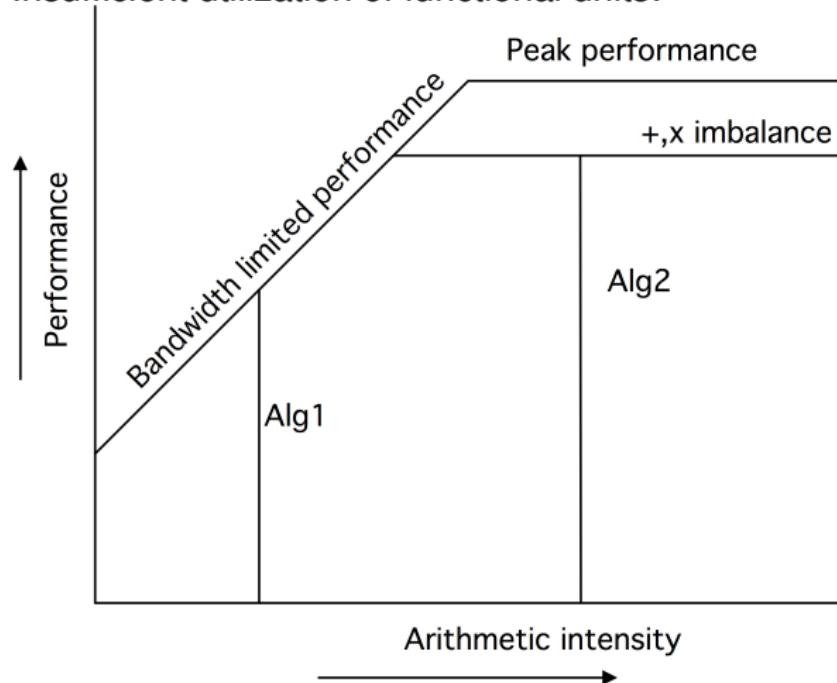
Roofline plot

Performance depends on algorithm:




Lowering the roof, 1

Insufficient utilization of functional units:



Lowering the roof, 2

Imperfect data transfer:



Spatial and temporal locality

Temporal locality: use an item, use it again but from cache
efficient because second transfer cheaper.

Spatial locality: use an item, then use one 'close to it'
(for instance from same cacheline)
efficient because item is already reachable even though not used
before.

Architecture aware programming

- ▶ Cache size: block loops
- ▶ pipelining and vector instructions: expose streams of instructions
- ▶ reuse: restructure code (both loop merge and splitting, unroll)
- ▶ TLB: don't jump all over memory
- ▶ associativity: watch out for powers of 2

Loop blocking

Multiple passes over data

```
for ( k< small bound )
    for ( i < N )
        x[i] = f( x[i], k, .... )
```

Block to be cache contained

```
for ( ii < N; ii+= blocksize )
    for ( k< small bound )
        for ( i=ii; i<ii+blocksize; i++ )
            x[i] = f( x[i], k, .... )
```

This requires independence of operations

The ultimate in performance programming: DGEMM

Matrix-matrix product $C = A \cdot B$

$$\forall_i \forall_j \forall_k : c_{ij} += a_{ik} b_{kj}$$

- ▶ Three independent loop i, j, k
- ▶ all three blocked i', j', k'
- ▶ Many loop permutations, blocking factors to choose

DGEMM variant

Inner products

```
for ( i )
    for ( j )
        for ( k )
            c[i, j] += a[i, k] * b[k, j]
```

DGEMM variant

Outer product: updates with low-rank columns-times-vector

```
for ( k )
    for ( i )
        for ( j )
            c[i,j] += a[i,k] * b[k,j]
```

DGEMM variant


Building up rows by linear combinations

```
for ( i )
    for ( k )
        for ( j )
            c[i, j] += a[i, k] * b[k, j]
```

Exchanging i, j : building up columns

Rank 1 updates

$$C_{**} = \sum_k A_{*k} B_{k*}$$



Matrix-panel multiply

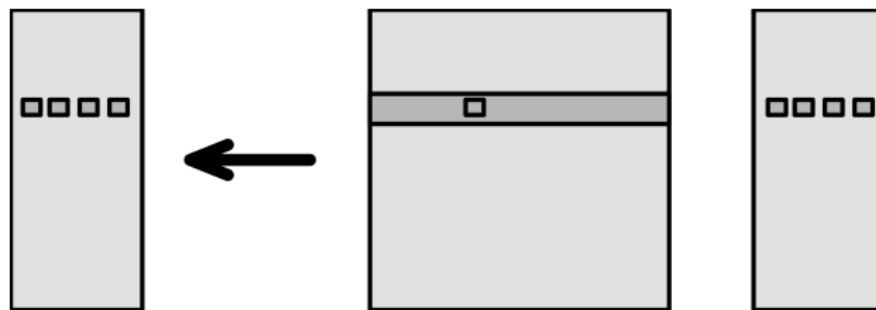
Block of A times ‘sliver’ of B



Inner algorithm

For inner i :

```
// compute C[i,*] :  
for k:  
    C[i,*] = A[i,k] * B[k,*]
```



Tuning




For inner i :

```
// compute C[i,*] :  
for k:  
    C[i,*] += A[i,k] * B[k,*]
```

- ▶ C[i,*] stays in register
- ▶ A[i,k] and B[k,*] stream from L1
- ▶ blocksize of A for L2 size
- ▶ A stored by rows to prevent TLB problems

Cache-oblivious programming

Observation: recursive subdivision will ultimately make a problem small / well-behaved enough



Cache-oblivious matrix-matrix multiply

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}$$

with $C_{11} = A_{11}B_{11} + A_{12}B_{21}$

Recursive approach will be cache contained.

Not as high performance as being cache-aware...