

COE 379

2025

Debugging is like being the detective in a crime movie where you are also the murderer. (Filipe Fortes, 2013)

When a program misbehaves, debugging is the process of finding out why. There are various strategies of finding errors in a program. The crudest one is debugging by print statements. If you have a notion of where in your code the error arises, you can edit your code to insert print statements, recompile, rerun, and see if the output gives you any suggestions. There are several problems with this:

- The edit/compile/run cycle is time consuming, especially since
- often the error will be caused by an earlier section of code, requiring you to edit, compile, and rerun repeatedly. Furthermore,
- the amount of data produced by your program can be too large to display and inspect effectively, and
- if your program is parallel, you probably need to print out data from all processors, making the inspection process very tedious.

For these reasons, the best way to debug is by the use of an interactive debugger, a program that allows you to monitor and control the behavior of a running program. In this section you will familiarize yourself with gdb and lldb, the open source debuggers of the GNU and clang projects respectively. Other debuggers are proprietary, and typically come with a compiler suite. Another distinction is that gdb is a commandline debugger; there are graphical debuggers such as ddd (a frontend to gdb) or DDT and TotalView (debuggers for parallel codes). We limit ourselves to gdb, since it incorporates the basic concepts common to all debuggers.

In this tutorial you will debug a number of simple programs with gdb and valgrind. The files can be found in the repository in the directory code/gdb.

## 1 Compiling for debug

You often need to recompile your code before you can debug it. A first reason for this is that the binary code typically knows nothing about what variable names corresponded to what memory locations, or what lines in the source to what instructions. In order to make the binary executable know this, you have to include the symbol table in it, which is done by adding the -g option to the compiler line.

Usually, you also need to lower the compiler optimization level: a production code will often be compiled with flags such as -O2 or -Xhost that try to make the code as fast as possible, but for debugging you need to replace this by -O0 ('oh-zero'). The reason is that higher levels will reorganize your code, making it hard to relate the execution to the source<sup>1</sup>.

---

1. Typically, actual code motion is done by -O3, but at level -O2 the compiler will inline functions and make other simplifications.

## 2 Invoking the debugger

There are three ways of using gdb: using it to start a program, attaching it to an already running program, or using it to inspect a core dump. We will only consider the first possibility.

| Starting a debugger run     |                               |
|-----------------------------|-------------------------------|
| gdb                         | lldb                          |
| \$ gdb program<br>(gdb) run | \$ lldb program<br>(lldb) run |

Here is an example of how to start gdb with program that has no arguments (Fortran users, use hello.F):

```
tutorials/gdb/c/hello.c
#include <stdlib.h>
#include <stdio.h>
int main() {
    printf("hello world\n");
    return 0;
}
%% cc -g -o hello hello.c
# regular invocation:
%% ./hello
hello world
# invocation from gdb:
%% gdb hello
GNU gdb 6.3.50-20050815 # .... [version info]
Copyright 2004 Free Software Foundation, Inc. .... [copyright info] ....
(gdb) run
Starting program: /home/eijkhout/tutorials/gdb/hello
Reading symbols for shared libraries +. done
hello world
```

Program exited normally.

```
(gdb) quit
%%
```

Important note: the program was compiled with the debug flag -g. This causes the symbol table (that is, the translation from machine address to program variables) and other debug information to be included in the binary. This will make your binary larger than strictly necessary, but it will also make it slower, for instance because the compiler will not perform certain optimizations<sup>2</sup>.

To illustrate the presence of the symbol table do

```
%% cc -g -o hello hello.c
%% gdb hello
GNU gdb 6.3.50-20050815 # .... version info
(gdb) list
```

and compare it with leaving out the -g flag:

---

2. Compiler optimizations are not supposed to change the semantics of a program, but sometimes do. This can lead to the nightmare scenario where a program crashes or gives incorrect results, but magically works correctly with compiled with debug and run in a debugger.

```

%% cc -o hello hello.c
%% gdb hello
GNU gdb 6.3.50-20050815 # .... version info
(gdb) list

```

For a program with commandline input we give the arguments to the run command (Fortran users use say.F):

|                       |  |
|-----------------------|--|
| tutorials/gdb/c/say.c | <pre> %% cc -o say -g say.c %% ./say 2 hello world hello world %% gdb say .... the usual messages ... (gdb) run 2 Starting program: /home/eijkhout/tutorials/gdb/c/say 2 Reading symbols for shared libraries +. done hello world hello world  Program exited normally. </pre> |
|-----------------------|--|

```

#include <stdlib.h>
#include <stdio.h>
int main(int argc,char **argv) {
    int i;
    for (i=0; i<atoi(argv[1]); i++)
        printf("hello world\n");
    return 0;
}

```

### 3 Finding errors: where, frame, print

Let us now consider some programs with errors.

#### 3.1 C programs

The following code has several errors. We will use the debugger to uncover them.

```

// /Users/eijkhout/Current/istc/scientific-computing-private/tutorials/gdb/c/square.c
int nmax,i;
float *squares,sum;

fscanf(stdin,"%d",nmax);
for (i=1; i<=nmax; i++) {
    squares[i] = 1.0/(i*i); sum += squares[i];
}
printf("Sum: %e\n",sum);

%% cc -g -o square square.c
%% ./square
5000
Segmentation fault

```

The segmentation fault (other messages are possible too) indicates that we are accessing memory that we are not allowed to, making the program exit. A debugger will quickly tell us where this happens:

```

%% gdb square
(gdb) run
50000

```

Program received signal EXC\_BAD\_ACCESS, Could not access memory.

Reason: KERN\_INVALID\_ADDRESS at address: 0x000000000000eb4a  
0x00007fff824295ca in \_\_svfscanf\_l ()

Apparently the error occurred in a function \_\_svfscanf\_l, which is not one of ours, but a system function. Using the backtrace (or bt, also where or w) command we display the call stack. This usually allows us to find out where the error lies:

Displaying a stack trace

---

gdb lldb

---

(gdb) where (lldb) thread backtrace

---

(gdb) where

```
#0 0x00007fff824295ca in __svfscanf_l ()  
#1 0x00007fff8244011b in fscanf ()  
#2 0x0000000100000e89 in main (argc=1, argv=0x7fff5fbfc7c0) at square.c:7
```

We inspect the actual problem:

---

Investigate a specific frame

---

gdb clang  
frame 2 frame select 2

---

We take a close look at line 7, and see that we need to change nmax to &nmax.

There is still an error in our program:

```
(gdb) run  
50000
```

Program received signal EXC\_BAD\_ACCESS, Could not access memory.  
Reason: KERN\_PROTECTION\_FAILURE at address: 0x000000010000f000  
0x00000000100000ebe in main (argc=2, argv=0x7fff5fbfc7a8) at square1.c:9  
9 squares[i] = 1.0/(i\*i); sum += squares[i];

We investigate further:

```
(gdb) print i  
$1 = 11237  
(gdb) print squares[i]  
Cannot access memory at address 0x10000f000  
(gdb) print squares  
$2 = (float *) 0x0
```

and we quickly see that we forgot to allocate squares.

By the way, we were lucky here: this sort of memory errors is not always detected. Starting our programm with a smaller input does not lead to an error:

```
(gdb) run  
50  
Sum: 1.625133e+00
```

Program exited normally.

Memory errors can also occur if we have a legitimate array, but we access it outside its bounds. The following program fills an array, forward, and reads it out, backward. However, there is an indexing error in the second loop.

```
// /Users/eijkhout/Current/istc/scientific-computing-private/tutorials/gdb/c/up.c  
int nlocal = 100,i;
```

```

double s, *array = (double*) malloc(nlocal*sizeof(double));
for (i=0; i<nlocal; i++) {
    double di = (double)i;
    array[i] = 1/(di*di);
}
s = 0.;
for (i=nlocal-1; i>=0; i++) {
    double di = (double)i;
    s += array[i];
}

Program received signal EXC_BAD_ACCESS, Could not access memory.
Reason: KERN_INVALID_ADDRESS at address: 0x0000000100200000
0x0000000100000f43 in main (argc=1, argv=0x7fff5fbfe2c0) at up.c:15
15      s += array[i];
(gdb) print array
$1 = (double *) 0x100104d00
(gdb) print i
$2 = 128608

```

You see that the index where the debugger finally complains is quite a bit larger than the size of the array.

Exercise 1. Can you think of a reason why indexing out of bounds is not immediately fatal? What would determine where it does become a problem? (Hint: how is computer memory structured?)

In section 8 you will see a tool that spots any out-of-bound indexing.

### 3.2 Fortran programs

Compile and run the following program:

```
MISSING SNIPPET gdb-squaref in codesnippetsdir=snippets.code
```

It should end prematurely with a message such as ‘Illegal instruction’. Running the program in gdb quickly tells you where the problem lies:

```
(gdb) run
Starting program: tutorials/gdb//fsquare
Reading symbols for shared libraries +++. done
```

```
Program received signal EXC_BAD_INSTRUCTION,
Illegal instruction/operand.
0x0000000100000da3 in square () at square.F:7
7      sum = sum + squares(i)
```

We take a close look at the code and see that we did not allocate squares properly.

## 4 Stepping through a program

---

### Stepping through a program

---

|      |      |                                    |
|------|------|------------------------------------|
| gdb  | lldb | meaning                            |
| run  |      | start a run                        |
| cont |      | continue from breakpoint           |
| next |      | next statement on same level       |
| step |      | next statement, this level or next |

---

Often the error in a program is sufficiently obscure that you need to investigate the program run in detail. Compile the following program

```
// /Users/eijkhout/Current/istc/scientific-computing-private/tutorials/gdb/c/roots.c
float root(int n)
{
    float r;
    r = sqrt(n);
    return r;
}

int main() {
    feenableexcept(FE_INVALID | FE_OVERFLOW);
    int i;
    float x=0;
    for (i=100; i>-100; i--)
        x += root(i+5);
    printf("sum: %e\n",x);
```

and run it:

```
%% ./roots
sum: nan
```

Start it in gdb as before:

```
%% gdb roots
GNU gdb 6.3.50-20050815
Copyright 2004 Free Software Foundation, Inc.
```

....

but before you run the program, you set a breakpoint at main. This tells the execution to stop, or ‘break’, in the main program.

```
(gdb) break main
Breakpoint 1 at 0x100000ea6: file root.c, line 14.
```

Now the program will stop at the first executable statement in main:

```
(gdb) run
Starting program: tutorials/gdb/c/roots
Reading symbols for shared libraries +. done
```

```
Breakpoint 1, main () at roots.c:14
14      float x=0;
```

Most of the time you will set a breakpoint at a specific line:

Set a breakpoint at a line

---

|                |                                   |
|----------------|-----------------------------------|
| gdb            | lldb                              |
| break foo.c:12 | breakpoint set [ -f foo.c ] -l 12 |

---

If execution is stopped at a breakpoint, you can do various things, such as issuing the step command:

```
Breakpoint 1, main () at roots.c:14
14      float x=0;
(gdb) step
15      for (i=100; i>-100; i--)
(gdb)
16      x += root(i);
(gdb)
```

(if you just hit return, the previously issued command is repeated). Do a number of steps in a row by hitting return. What do you notice about the function and the loop?

Switch from doing step to doing next. Now what do you notice about the loop and the function?

Set another breakpoint: break 17 and do cont. What happens?

Rerun the program after you set a breakpoint on the line with the sqrt call. When the execution stops there do where and list.

- If you set many breakpoints, you can find out what they are with info breakpoints.
- You can remove breakpoints with delete n where n is the number of the breakpoint.
- If you restart your program with run without leaving gdb, the breakpoints stay in effect.
- If you leave gdb, the breakpoints are cleared but you can save them: save breakpoints <file>. Use source <file> to read them in on the next gdb run.

## 5 Inspecting values

Run the previous program again in gdb: set a breakpoint at the line that does the sqrt call before you actually call run. When the program gets to line 8 you can do print n. Do cont. Where does the program stop?

If you want to repair a variable, you can do set var=value. Change the variable n and confirm that the square root of the new value is computed. Which commands do you do?

## 6 Breakpoints

If a problem occurs in a loop, it can be tedious keep typing cont and inspecting the variable with print. Instead you can add a condition to an existing breakpoint. First of all, you can make the breakpoint subject to a condition: with

condition 1 if (n<0)

breakpoint 1 will only obeyed if n<0 is true.

You can also have a breakpoint that is only activated by some condition. The statement

break 8 if (n<0)

means that breakpoint 8 becomes (unconditionally) active after the condition n<0 is encountered.

Set a breakpoint

| gdb                   | lldb                              |
|-----------------------|-----------------------------------|
| break foo.c:12        | breakpoint set [ -f foo.c ] -l 12 |
| break foo.c:12 if n>0 |                                   |

Remark You can break on NaN with the following trick:

break foo.c:12 if x!=x

using the fact that NaN is the only number not equal to itself.

Another possibility is to use ignore 1 50, which will not stop at breakpoint 1 the next 50 times.

Remove the existing breakpoint, redefine it with the condition n<0 and rerun your program. When the program breaks, find for what value of the loop variable it happened. What is the sequence of commands you use?

You can set a breakpoint in various ways:

- break foo.c to stop when code in a certain file is reached;
- break 123 to stop at a certain line in the current file;
- break foo to stop at subprogram foo
- or various combinations, such as break foo.c:123.

Information about breakpoints:

- If you set many breakpoints, you can find out what they are with info breakpoints.
- You can remove breakpoints with delete n where n is the number of the breakpoint.
- If you restart your program with run without leaving gdb, the breakpoints stay in effect.
- If you leave gdb, the breakpoints are cleared but you can save them: save breakpoints <file>. Use source <file> to read them in on the next gdb run.
- In languages with exceptions, such as C++, you can set a catchpoint:

Set a breakpoint for exceptions

---

|             |                  |
|-------------|------------------|
| gdb         | clang            |
| catch throw | break set -E C++ |

---

Finally, you can execute commands at a breakpoint:

```
break 45
command
print x
cont
end
```

This states that at line 45 variable x is to be printed, and execution should immediately continue.

If you want to run repeated gdb sessions on the same program, you may want to save an reload breakpoints. This can be done with

```
save-breakpoint filename
source filename
```

## 7 Memory debugging

Many problems in programming stem from memory errors. We start with a sort description of the most common types, and then discuss tools that help you detect them.

### 7.1 Type of memory errors

#### 7.1.1 Invalid pointers

Dereferencing a pointer that does not point to an allocated object can lead to an error. If your pointer points into valid memory anyway, your computation will continue but with incorrect results.

However, it is more likely that your program will probably exit with a segmentation violation or a bus error.

#### 7.1.2 Out-of-bounds errors

Addressing outside the bounds of an allocated object is less likely to crash your program and more likely to give incorrect results.

Exceeding bounds by a large enough amount will again give a segmentation violation, but going out of bounds by a small amount may read invalid data, or corrupt data of other variables, giving incorrect results that may go undetected for a long time.

### 7.1.3 Memory leaks

We speak of a memory leak if allocated memory becomes unreachable. Example:

```
if (something) {
    double *x = malloc(10*sizeof(double));
    // do something with x
}
```

After the conditional, the allocated memory is not freed, but the pointer that pointed to has gone away.

This last type especially can be hard to find. Memory leaks will only surface in that your program runs out of memory. That in turn is detectable because your allocation will fail. It is a good idea to always check the return result of your malloc or allocate statement!

## 8 Memory debugging with Valgrind

Errors leading to memory problems are easy to make. In this section we will see how valgrind makes it possible to track down these errors. The use of valgrind is simplicity itself:

```
valgrind yourprogram yourargs
```

As a first example, consider out of bound addressing, also known as buffer overflow:

```
// corruption/cxx/outofbound.cpp
vector<float> x(n);
cout << x[n] << '\n';
```

This is unlikely to crash your code, but the results are unpredictable, and this is certainly a failure of your program logic.

Valgrind indicates that this is an invalid read, what line it occurs on, and where the block was allocated:

```
==9112== Invalid read of size 4
==9112==   at 0x40233B: main (outofbound.cpp:10)
==9112== Address 0x595fde8 is 0 bytes after a block of size 40 alloc'd
==9112==   at 0x4C2A483: operator new(unsigned long) (vg_replace_malloc.c:344)
==9112==   by 0x4023CD: allocate (new_allocator.h:111)
==9112==   by 0x4023CD: allocate (alloc_traits.h:436)
==9112==   by 0x4023CD: __M_allocate (stl_vector.h:296)
==9112==   by 0x4023CD: __M_create_storage (stl_vector.h:311)
==9112==   by 0x4023CD: __Vector_base (stl_vector.h:260)
==9112==   by 0x4023CD: __Vector_base (stl_vector.h:258)
==9112==   by 0x4023CD: vector (stl_vector.h:415)
==9112==   by 0x4023CD: main (outofbound.cpp:9)
```

Remark Buffer overflows are a well-known security risk, typically associated with reading string input from a user source. Buffer overflows can be largely avoided by using C++ constructs such as `cin` and `string` instead of `sscanf` and character arrays.

Valgrind is informative but cryptic, since it works on the bare memory, not on variables. Thus, these error messages take some exegesis. They state that line 10 reads a 4-byte object immediately after a block of 40 bytes that was allocated. In other words: the code is writing outside the bounds of an allocated array. The output also gives an indication where in the source that allocation was performed, which is usually enough to pinpoint the problem.

The next example performs a read on an array that has already been free'd. In this simple case you will actually get the expected output, but if the read comes much later than the free, the output can be anything. (Do you understand how that can happen?)

```
// corruption/cxx/free.cpp
auto ar = new float[10000];
ar[n] = 3.14;
delete [] ar;
cout << ar[n] << '\n';
```

Valgrind again states that this is an invalid read; it gives both where the block was allocated and where it was freed.

```
==11431== Invalid read of size 4
==11431==   at 0x40230C: main (free.cpp:13)
==11431== Address 0x595fdcc is 12 bytes inside a block of size 40,000 free'd
==11431==   at 0x4C2B5CF: operator delete(void*, unsigned long) (vg_replace_malloc.c:595)
==11431== by 0x402301: main (free.cpp:12)
==11431== Block was alloc'd at
==11431==   at 0x4C2AB28: operator new[](unsigned long) (vg_replace_malloc.c:433)
==11431== by 0x4022E0: main (free.cpp:10)
```

On the other hand, if you forget to free memory you have a memory leak (just imagine allocating, and not free'ing, in a loop; section ??).

```
// corruption/cxx/leak.cpp
auto ar = new float[10000];
ar[n] = 3.14;
cout << ar[n] << '\n';
return 0;
```

which valgrind reports on:

```
==283234== LEAK SUMMARY:
==283234==   definitely lost: 40,000 bytes in 1 blocks
==283234==   indirectly lost: 0 bytes in 0 blocks
==283234==   possibly lost: 0 bytes in 0 blocks
==283234==   still reachable: 8 bytes in 1 blocks
==283234==   suppressed: 0 bytes in 0 blocks
```

Memory leaks are much more rare in C++ than in C because of containers such as `std::vector`. However, in sophisticated cases you may still do your own memory management, and you need to be aware of the danger of memory leaks.

If you do your own memory management, there is also a danger of writing to an array pointer that has not been allocated yet:

```
// corruption/cxx/init.cpp
float* ar;
ar[n] = 3.14;
cout << ar[n] << '\n';
```

The behavior of this code depends on all sorts of things: if the pointer variable is zero, the code will crash. On the other hand, if it contains some random value, the write may succeed; provided you are not writing too far from that location.

The output here shows both the valgrind diagnosis, and the Operating System (OS) message when the program aborted:

```
==283234== LEAK SUMMARY:
==283234==   definitely lost: 40,000 bytes in 1 blocks
==283234==   indirectly lost: 0 bytes in 0 blocks
==283234==   possibly lost: 0 bytes in 0 blocks
==283234==   still reachable: 8 bytes in 1 blocks
==283234==   suppressed: 0 bytes in 0 blocks
```

### 8.1 Electric fence

The electric fence library is one of a number of tools that supplies a new malloc with debugging support. These are linked instead of the malloc of the standard libc.

```
cc -o program program.c -L/location/of/efence -lefence
```

Suppose your program has an out-of-bounds error. Running with gdb, this error may only become apparent if the bounds are exceeded by a large amount. On the other hand, if the code is linked with libefence, the debugger will stop at the very first time the bounds are exceeded.

## 9 Further reading

A good tutorial: <http://www.dirac.org/linux/gdb/>.

Reference manual: [http://www.ofb.net/gnu/gdb/gdb\\_toc.html](http://www.ofb.net/gnu/gdb/gdb_toc.html).

|                                       |   |                                  |    |
|---------------------------------------|---|----------------------------------|----|
| Contents                              | 5 | Inspecting values                | 7  |
| 1 Compiling for debug                 | 1 | 6 Breakpoints                    | 7  |
| 2 Invoking the debugger               | 2 | 7 Memory debugging               | 8  |
| 3 Finding errors: where, frame, print | 3 | 7.1 Type of memory errors        | 8  |
| 3.1 C programs                        | 3 | 8 Memory debugging with Valgrind | 9  |
| 3.2 Fortran programs                  | 5 | 8.1 Electric fence               | 11 |
| 4 Stepping through a program          | 5 | 9 Further reading                | 11 |

## References

- [1] John L. Hennessy and David A. Patterson. Computer Architecture, A Quantitative Approach. Morgan Kaufman Publishers, 3rd edition edition, 1990, 3rd edition 2003. [Not cited.]

Table 1: List of common gdb / lldb commands.

| gdb   | lldb                               |
|---|------------------------------------|
| Starting a debugger run                                     |                                    |
| \$ gdb program<br>(gdb) run                                 | \$ lldb program<br>(lldb) run      |
| Displaying a stack trace                                    |                                    |
| (gdb) where   | (lldb) thread backtrace            |
| Investigate a specific frame                                |                                    |
| frame 2   | frame select 2                     |
| Run/step  |                                    |
| run / step / continue                                       | thread continue / step-in/over/out |
| Set a breakpoint at a line                                  |                                    |
| break foo.c:12<br>break foo.c:12 if n>0<br>info breakpoints | breakpoint set [ -f foo.c ] -l 12  |
| Set a breakpoint for exceptions                             |                                    |
| catch throw   | break set -E C++                   |