

# What is parallelism?

Victor Eijkhout

Fall 2025

## Justification

Parallel computing has been a necessity for decades in computational science. Here we discuss some of the basic concepts. Actual parallel programming will be discussed in other lectures.

# **Parallelism**

## **Basic concepts**

## The basic idea

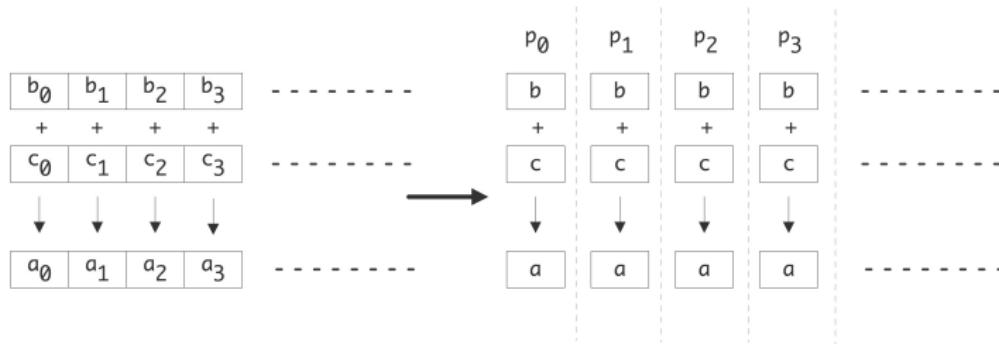
Parallelism is about doing multiple things at once.

- ▶ Hardware: vector instructions, multiple cores, nodes in a cluster.
- ▶ Algorithm: can you think of examples?

# Simple example

Summing two arrays together:

```
1 for (i=0; i<n; i++)  
2   a[i] = b[i] + c[i];
```



Parallel: every processing element does

```
1 for ( i in my_subset_of_indices )  
2   a[i] = b[i] + c[i];
```

Time goes down linearly with processors

## Differences between operations

```
1 for (i=0; i<n; i++)  
2   a[i] = b[i] + c[i];
```

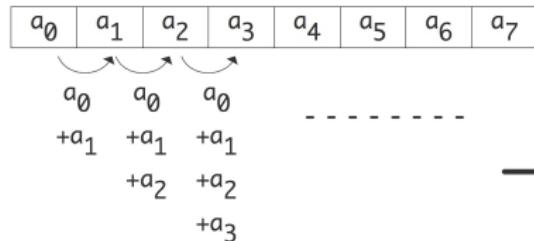
```
1 s = 0;  
2 for (i=0; i<n; i++)  
3   s += x[i]
```

- ▶ Compare operation counts
- ▶ Compare behavior on single processor. What about multi-core?
- ▶ Other thoughts about parallel execution?

# Summing

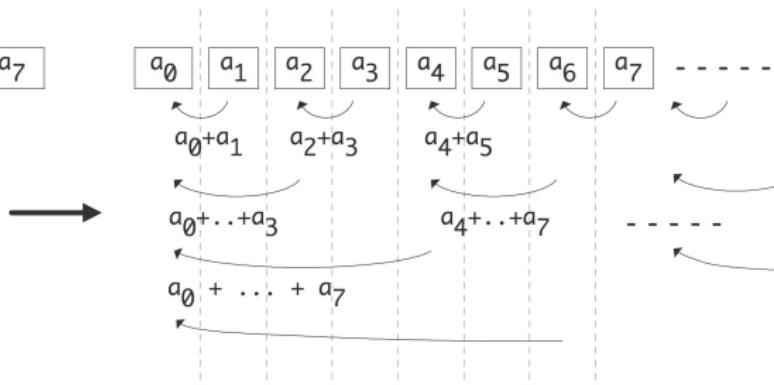
## Naive algorithm

```
1 s = 0;  
2 for (i=0; i<n; i++)  
3     s += x[i]
```



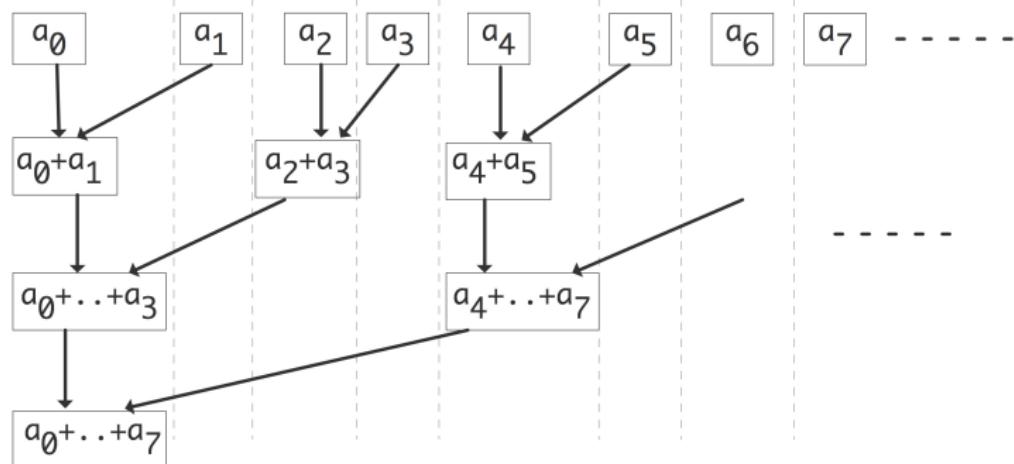
## Recoding

```
1 for (s=2; s<n; s*=2)  
2     for (i=0; i<n; i+=s)  
3         x[i] += x[i+s/2]
```



## And then there is hardware

Topology of the processors:



increasing distance: limit on parallel speedup

## **Theoretical concepts**

## **Efficiency and scaling**

## Speedup

- ▶ Single processor time  $T_1$ , on  $p$  processors  $T_p$
- ▶ speedup is  $S_p = T_1 / T_p$ ,  $S_p \leq p$
- ▶ efficiency is  $E_p = S_p / p$ ,  $0 < E_p \leq 1$

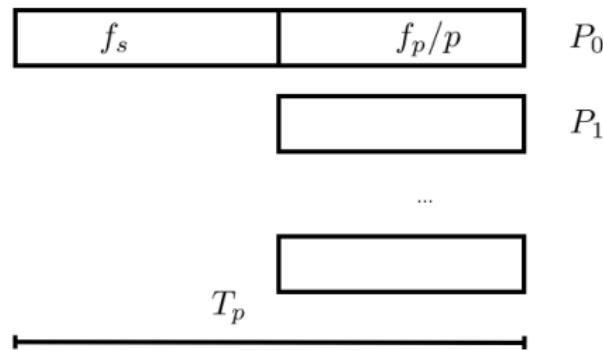
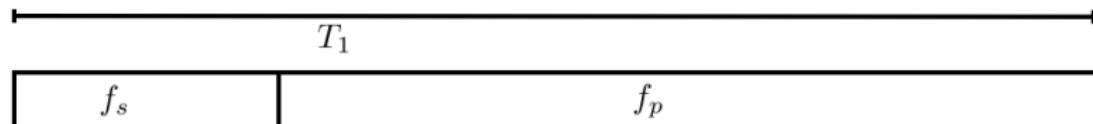
But:

- ▶ Is  $T_1$  based on the same algorithm? The parallel code?
- ▶ Sometimes superlinear speedup.
- ▶ Is  $T_1$  measurable? Can the problem be run on a single processor?

## Amdahl's law

Let's assume that part of the application can be parallelized, part not.  
(Examples?)

- ▶  $F_s$  sequential fraction,  $F_p$  parallelizable fraction
- ▶  $F_s + F_p = 1$



## Amdahl's law, analysis

- ▶  $F_s$  sequential fraction,  $F_p$  parallelizable fraction
- ▶  $F_s + F_p = 1$
- ▶  $T_1 = (F_s + F_p)T_1 = F_s T_1 + F_p T_1$
- ▶ Amdahl's law:  $T_p = F_s T_1 + F_p T_1 / p$
- ▶  $P \rightarrow \infty: T_P \downarrow T_1 F_s$
- ▶ Speedup is limited by  $S_P \leq 1/F_s$ , efficiency is a decreasing function  $E \sim 1/P$ .

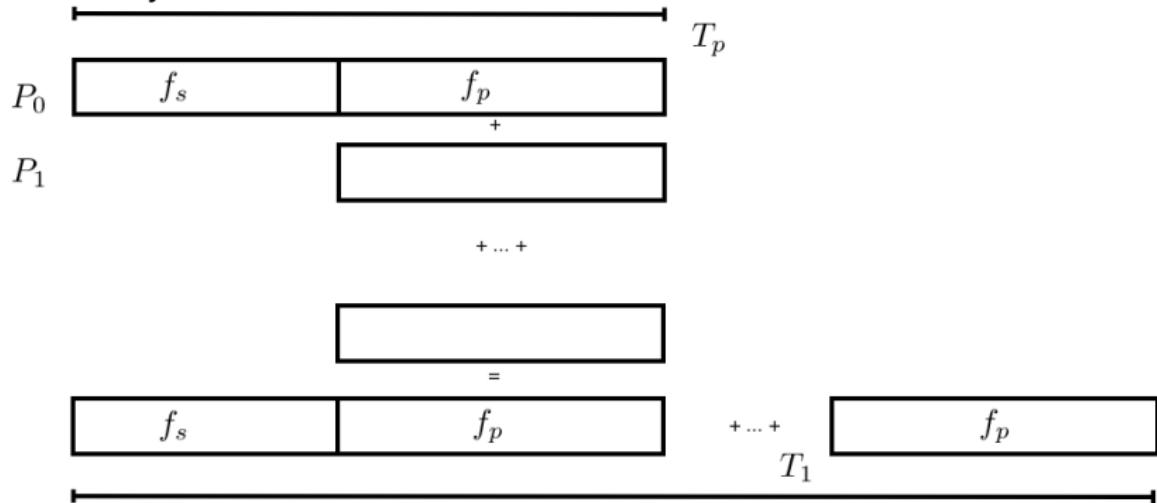
Do you see problems with this?

## Amdahl's law with communication overhead

- ▶ Communication independent of  $p$ :  $T_p = T_1(F_s + F_p/P) + T_c$
- ▶ assume fully parallelizable:  $F_p = 1$
- ▶ then  $S_p = \frac{T_1}{T_1/p + T_c}$
- ▶ For reasonable speedup:  $T_c \ll T_1/p$  or  $p \ll T_1/T_c$ :  
number of processors limited by ratio of scalar execution time and communication overhead

## Gustafson's law

Reconstruct the sequential execution from the parallel, then analyze efficiency.



## Gustafson's law

- ▶ Let  $T_p = F_s + F_p \equiv 1$
- ▶ then  $T_1 = F_s + p \cdot F_p$
- ▶ Speedup:

$$S_p = \frac{T_1}{T_p} = \frac{F_s + p \cdot F_p}{F_s + F_p} = F_s + p \cdot F_p = p - (p-1) \cdot F_s.$$

slowly decreasing function of  $p$

## Scaling

- ▶ Amdahl's law: strong scaling  
same problem over increasing processors
- ▶ Often more realistic: weak scaling  
increase problem size with number of processors,  
for instance keeping memory constant
- ▶ Weak scaling:  $E_p > c$
- ▶ example (below): dense linear algebra

## Strong scaling

- ▶ Let  $M$  be the total memory needed for your problem.
- ▶ Let  $P$  be the number of processors  
⇒ memory per processor is  $M/P$
- ▶ What is  $\lim_{P \rightarrow \infty} E_P$ ?  
(Note that implicitly  $E_p = E(P, M)$ .)

## Weak scaling

- ▶ Let  $M$  be the memory per processor.
- ▶ Let  $P$  be the number of processors  
⇒ total memory is  $M \cdot P$
- ▶ What is  $\lim_{P \rightarrow \infty} E_P$ ?  
(Note that implicitly  $E_p = E(P, M)$ .)

## Simulation scaling

- ▶ Assumption: simulated time  $S$ , running time  $T$  constant, now increase precision
- ▶  $m$  memory per processor, and  $P$  the number of processors

$$M = Pm \quad \text{total memory.}$$

$d$  the number of space dimensions of the problem, typically 2 or 3,

$$\Delta x = 1/M^{1/d} \quad \text{grid spacing.}$$

- ▶ stability:

$$\Delta t = \begin{cases} \Delta x = 1 / M^{1/d} & \text{hyperbolic case} \\ \Delta x^2 = 1 / M^{2/d} & \text{parabolic case} \end{cases}$$

With a simulated time  $S$ :

$$k = S/\Delta t \quad \text{time steps.}$$

## Simulation scaling con'td

- ▶ Assume time steps parallelizable

$$T = kM/P = \frac{S}{\Delta t} m.$$

Setting  $T/S = C$ , we find

$$m = C\Delta t,$$

memory per processor goes down.

$$m = C\Delta t = c \begin{cases} 1 / M^{1/d} & \text{hyperbolic case} \\ 1 / M^{2/d} & \text{parabolic case} \end{cases}$$

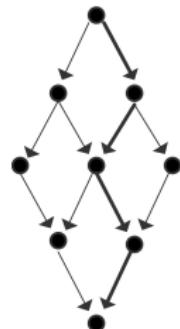
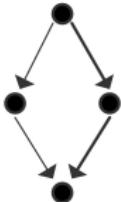
- ▶ Substituting  $M = Pm$ , we find ultimately

$$m = C \begin{cases} 1 / P^{1/(d+1)} & \text{hyperbolic} \\ 1 / P^{2/(d+2)} & \text{parabolic} \end{cases}$$

## **Critical path analysis**

## Critical path

- ▶ The sequential fraction contains a *critical path*: a sequence of operations that depend on each other.
- ▶ Example?
- ▶  $T_\infty$  = time with unlimited processors: length of critical path.



## Brent's theorem

Let  $m$  be the total number of tasks,  $p$  the number of processors, and  $t$  the length of a critical path. Then the computation can be done in

$$T_p \leq t + \frac{m-t}{p}.$$

- ▶ Time equals the length of the critical path ...
- ▶ ... plus the remaining work as parallel as possible.

# **Granularity**

## Definition

Definition: granularity is the measure for how many operations can be performed between synchronizations

## Instruction level parallelism

$$\begin{aligned}a &\leftarrow b + c \\d &\leftarrow e * f\end{aligned}$$

For the compiler / processor to worry about

## Data parallelism

```
1 for (i=0; i<1000000; i++)  
2   a[i] = 2*b[i];
```

- ▶ Array processors, vector instructions, pipelining, GPUs
- ▶ Sometimes harder to discover
- ▶ Often used mixed with other forms of parallelism

## Task-level parallelism

Unsynchronized tasks: fork-join  
general scheduler

## Conveniently parallel

Example: Mandelbrot set

Parameter sweep,  
often best handled by external tools

## Medium-grain parallelism

### Mix of data parallel and task parallel

```
1 my_lower_bound = // some processor-dependent number
2 my_upper_bound = // some processor-dependent number
3 for (i=my_lower_bound; i<my_upper_bound; i++)
4   // the loop body goes here
```

## **LU factorization analysis**

## Algorithm

```
for k = 1, n - 1:  
    for i = k + 1 to n:  
         $a_{ik} \leftarrow a_{ik} / a_{kk}$   
    for i = k + 1 to n:  
        for j = k + 1 to n:  
             $a_{ij} \leftarrow a_{ij} - a_{ik} * a_{kj}$ 
```

Can the  $k$  loop be done in parallel? The  $i, j$  loops?

## Dependent operations

$$a_{22} \leftarrow a_{22} - a_{21} * a_{11}^{-1} a_{12}$$

...

$$a_{33} \leftarrow a_{33} - a_{32} * a_{22}^{-1} a_{23}$$

## Exercise 1

Critical path Follow this argument through. Argue that there is a non-trivial critical path in the sense of section ?? . What is its length? In the analysis of the critical path section, what does this critical path imply for the minimum parallel execution time and bounds on speedup?

## Subblock update

for  $i = k + 1$  to  $n$ :

    for  $j = k + 1$  to  $n$ :

$$a_{ij} \leftarrow a_{ij} - a_{ik} * a_{kj}$$

How many processors can you use maximally in step  $k$ ?

## Exercise 2

Parallel execution Continue this reasoning. With  $p = n^2$  processing elements each of the  $(i,j)$  updates in the subblock can be done simultaneously. To be precise, how long does an arbitrary  $k$  iteration take? Summing over all  $k$ , what is the resulting  $T_p, S_p, E_p$ ? How does this relate to the bounds you derived above?

Also, with  $p = n$  processing elements you could let each row or column of the subblock update be done in parallel. What is now the time for the  $k$ th outer iteration? What is the resulting  $T_p, S_p, E_p$ ?,

## Application scaling

Single processor.

Relating time and memory to problem size

$$T = \frac{1}{3}N^3/f, \quad M = N^2.$$

where  $f$  is processor frequency.

## Exercise 3

Memory scaling, case 1: Faster processor Suppose you buy a processor twice as fast, and you want to do a benchmark run that again takes time  $T$ . How much memory do you need?

## More processors

Keep frequency constant, but vary number of processors  $p$ :

$$T = \frac{1}{3}N^3/p, \quad M = N^2.$$

Each processor now stores  $M_p = N^2/p$  elements.

## Exercise 4

Memory scaling, case 2: More processors Suppose you have a cluster with  $p$  processors, each with  $M_p$  memory, can run a Gaussian elimination of an  $N \times N$  matrix in time  $T$ :

$$T = \frac{1}{3}N^3/p, \quad M_p = N^2/p.$$

Now you extend the cluster to  $2P$  processors, of the same clock speed, and you want to do a benchmark run, again taking time  $T$ . How much memory does each node need?

Hint: for the extended cluster:

$$T' = \frac{1}{3}N'^3/p', \quad M'_p = N'^2/p'.$$

The question becomes to compute  $M'_p$  under the given conditions.

## **The SIMD/MIMD/SPMD/SIMT model for parallelism**

## Flynn Taxonomy

Consider instruction stream and data stream:

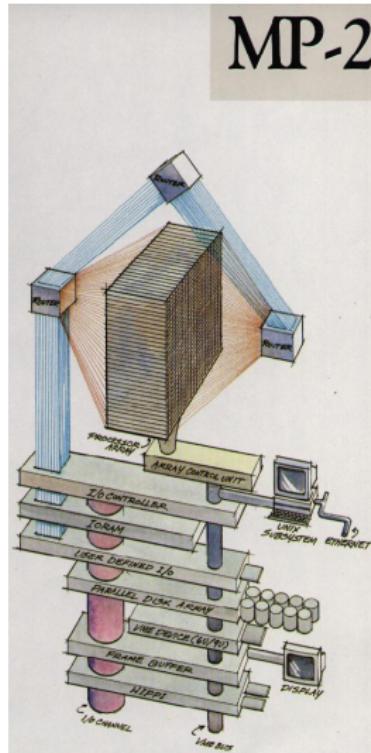
- ▶ SISD: single instruction single data  
used to be single processor, now single core
- ▶ MISD: multiple instruction single data  
redundant computing for fault tolerance?
- ▶ SIMD: single instruction multiple data  
data parallelism, pipelining, array processing, vector instructions
- ▶ MIMD: multiple instruction multiple data  
independent processors, clusters, MPPs

## SIMD

- ▶ Relies on streams of identical operations
- ▶ See pipelining
- ▶ Recurrences hard to accomodate

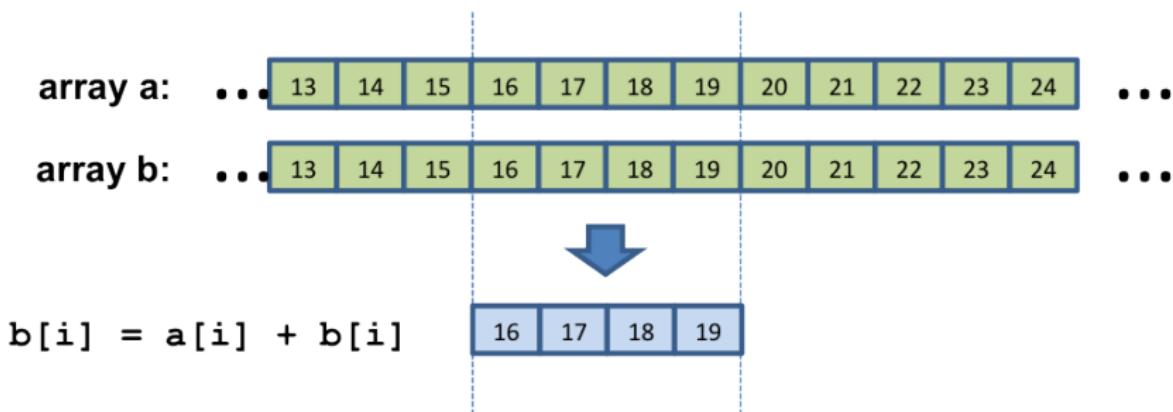
# SIMD: array processors

Technology going back to the 1980s: FPS, MasPar, CM, GoodYear  
Major advantage: simplification of processor



# SIMD as vector instructions

- ▶ Register width multiple of 8 bytes:
- ▶ simultaneous processing of more than one operand pair
- ▶ SSE: 2 operands,
- ▶ AVX: 4 or 8 operands



# Controlling vector instructions

```
1 void func(float *restrict c, float *restrict a,
2           float *restrict b, int n)
3 {
4     #pragma vector always
5     for (int i=0; i<n; i++)
6         c[i] = a[i] * b[i];
7 }
```

This needs aligned data (posix\_memalign)

## New branches in the taxonomy

- ▶ SPMD: single program multiple data  
the way clusters are actually used
- ▶ SIMD: single instruction multiple threads  
the GPU model

## MIMD becomes SPMD

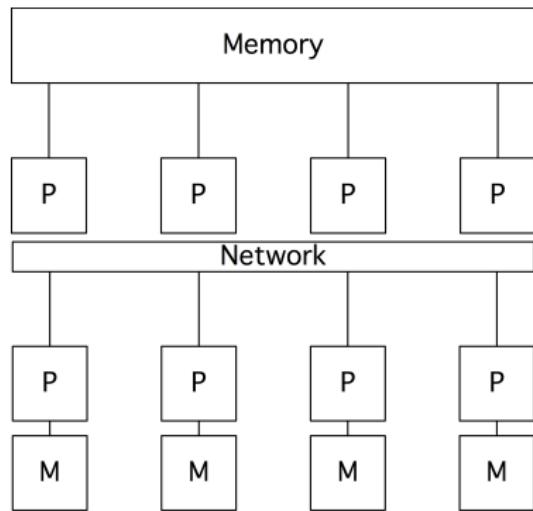
- ▶ MIMD: independent processors, independent instruction streams, independent data
- ▶ In practice very little true independence: usually the same executable  
Single Program Multiple Data
- ▶ Exceptional example: climate codes
- ▶ Old-style SPMD: cluster of single-processor nodes
- ▶ New-style: cluster of multicore nodes, ignore shared caches / memory
- ▶ (We'll get to hybrid computing in a minute)

## GPUs and data parallelism

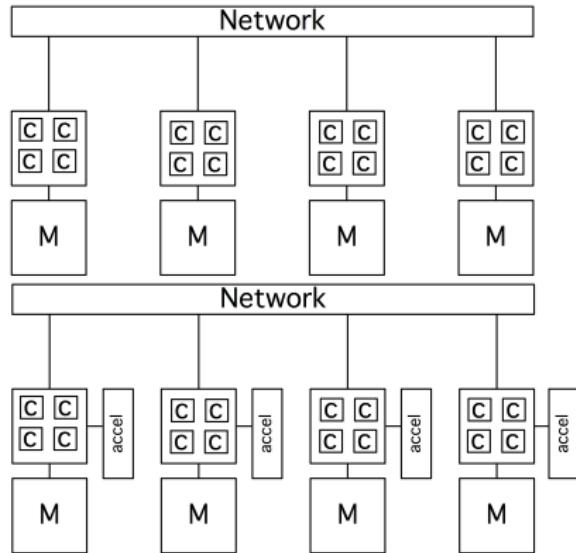
Lockstep in thread block,  
single instruction model between streaming processors  
(more about GPU threads later)

## **Characterization of parallelism by memory model**

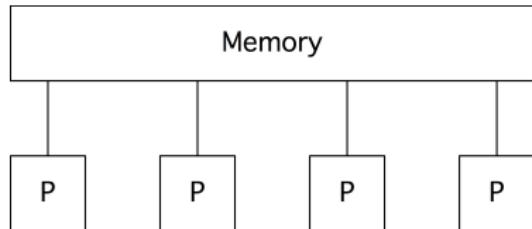
## Major types of memory organization, classic



# Major types of memory organization, contemporary



## Symmetric multi-processing



- ▶ The ideal case of shared memory:  
every address equally accessible
- ▶ This hasn't existed in a while  
(Tim Mattson claims Cray-2)
- ▶ Danger signs: shared memory programming pretends that  
memory access is symmetric  
in fact: hides reality from you

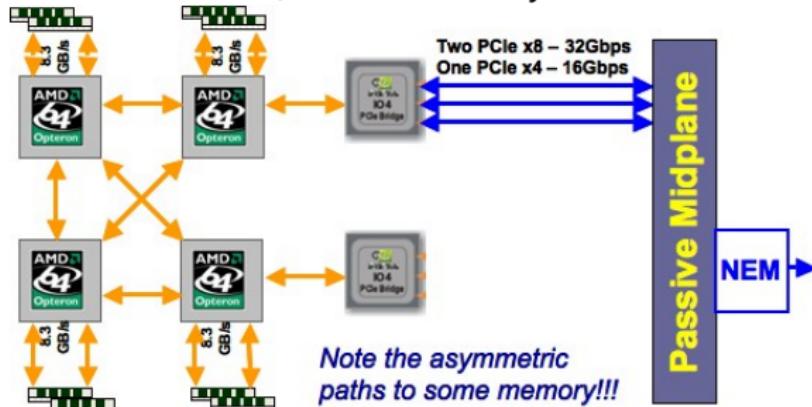
## SMP, bus design

- ▶ Bus: all processors on the same wires to memory
- ▶ Not very scalable: requires slow processors or cache memory
- ▶ Cache coherence easy by ‘snooping’

# Non-uniform Memory Access

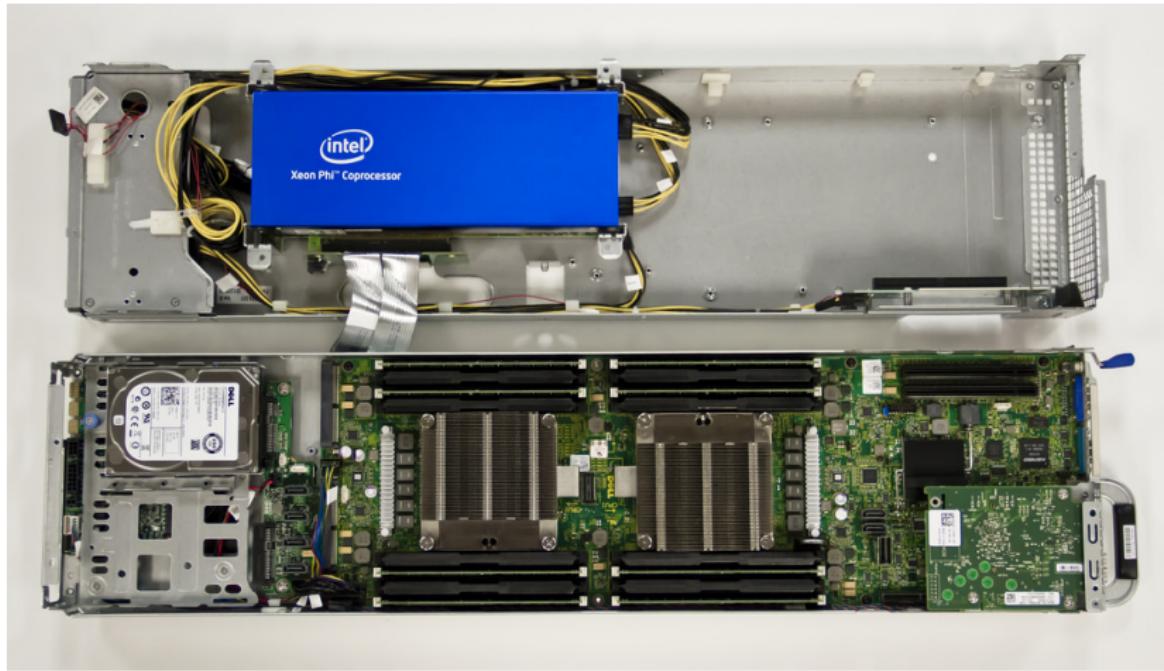
Memory is equally programmable, but not equally accessible

- ▶ Different caches, different affinity



- ▶ Distributed shared memory: network latency  
ScaleMP and other products watch me not believe it

# Picture of NUMA



## **Interconnects and topologies, theoretical concepts**

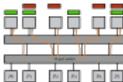
## Topology concepts

- ▶ Hardware characteristics
- ▶ Software requirement
- ▶ Design: how 'close' are processors?

## Graph theory

- ▶ Degree: number of connections from one processor to others
- ▶ Diameter: maximum minimum distance (measured in hops)

# Bandwidth

- ▶ Bandwidth per wire is nice, adding over all wires is nice, but . . .  
A diagram illustrating a network connection. At the top, there are four small square icons representing hosts or switches. Four vertical lines descend from these icons to a horizontal row of four rectangular boxes. From each of these boxes, two more vertical lines descend to a second horizontal row of four rectangular boxes. This second row of boxes has two vertical lines extending downwards, representing the final connection to another set of hosts or switches.
- ▶ Bisection width: minimum number of wires through a cut
- ▶ Bisection bandwidth: bandwidth through a bisection

## Design 1: bus

Already discussed; simple design, does not scale very far

## Design 2: linear arrays

- ▶ Degree 2, diameter  $P$ , bisection width 1
- ▶ Scales nicely!
- ▶ but low bisection width

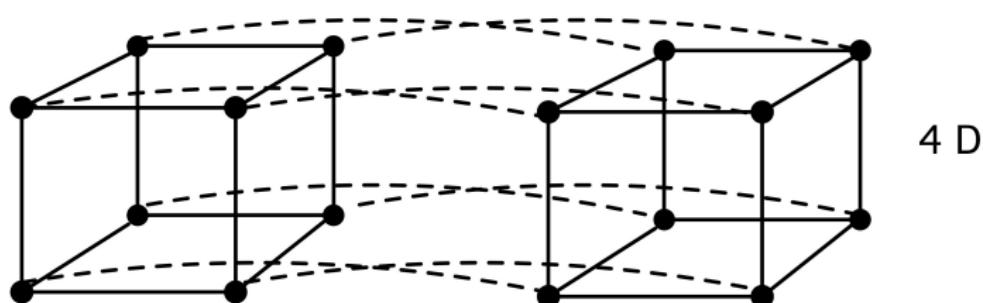
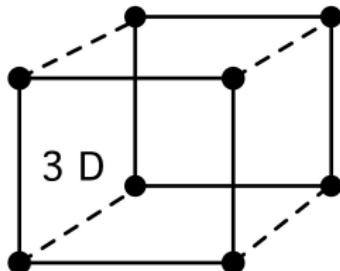
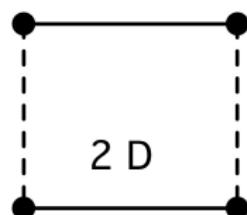
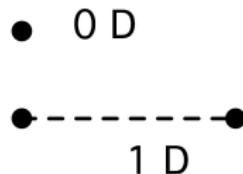
## Exercise 5

Broadcast algorithm Flip last bit, flip one before, . . .

## Design 3: 2/3-D arrays

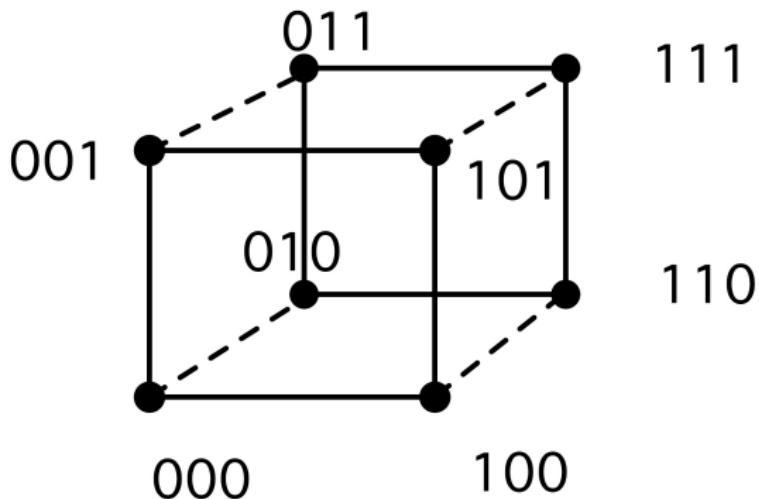
- ▶ Degree  $2d$ , diameter  $P^{1/d}$
- ▶ Natural design: nature is three-dimensional
- ▶ More dimensions: less contention.  
K-machine is 6-dimensional

## Design 3: Hypercubes



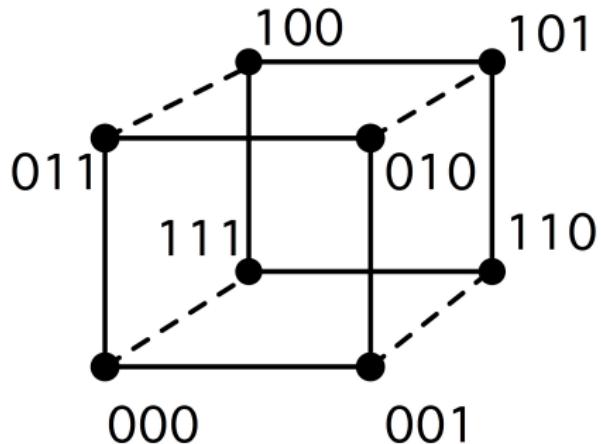
## Hypercube numbering

Naive numbering:



## Gray codes

Embedding linear numbering in hypercube:



## Binary reflected Gray code

1D Gray code :

0 1

2D Gray code :

1D code and reflection: 0 1 : 1 0

append 0 and 1 bit: 0 0 : 1 1

2D code and reflection: 0 1 1 0 : 0 1 1 0

3D Gray code :

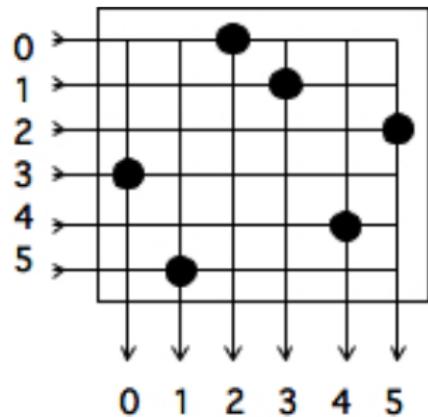
0 0 1 1 : 1 1 0 0

append 0 and 1 bit: 0 0 0 0 : 1 1 1 1

## Switching networks

- ▶ Solution to all-to-all connection
- ▶ (Real all-to-all too expensive)
- ▶ Typically layered
- ▶ Switching elements: easy to extend

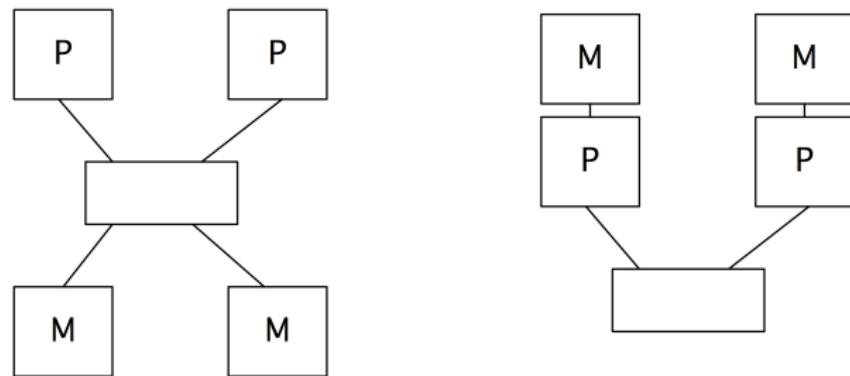
## Cross bar



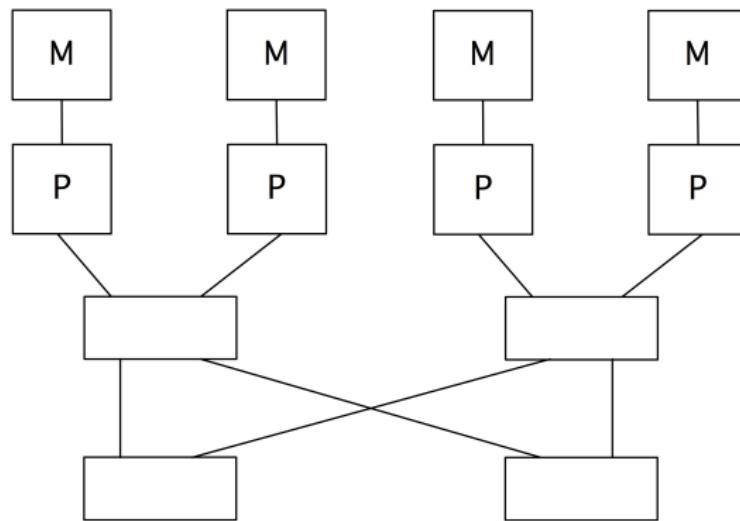
Advantage: non-blocking  
Disadvantage: cost

## Butterfly exchange

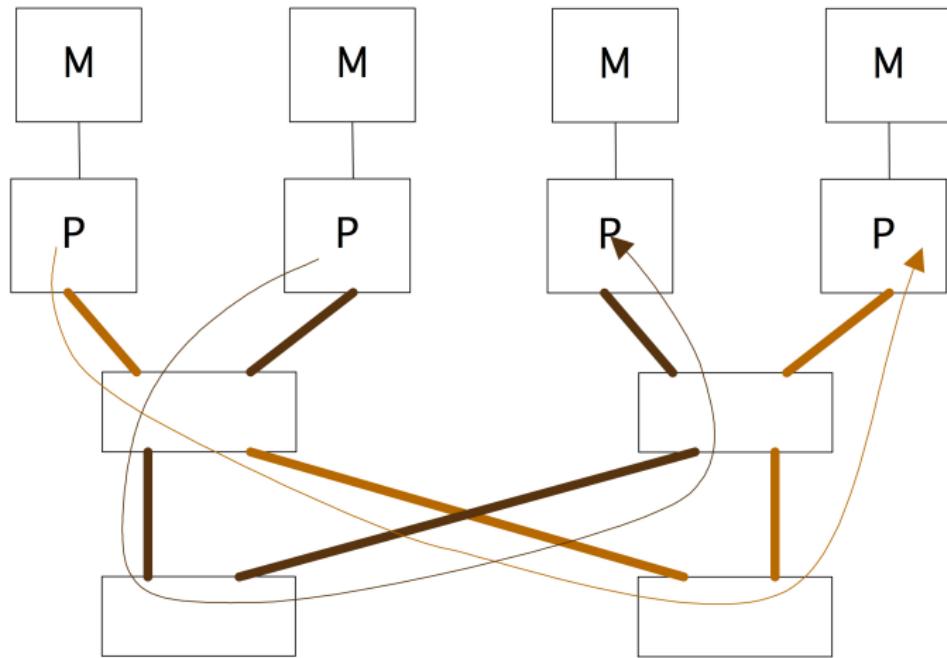
Process to segmented pool of memory, or between processors with private memory:



# Building up butterflies

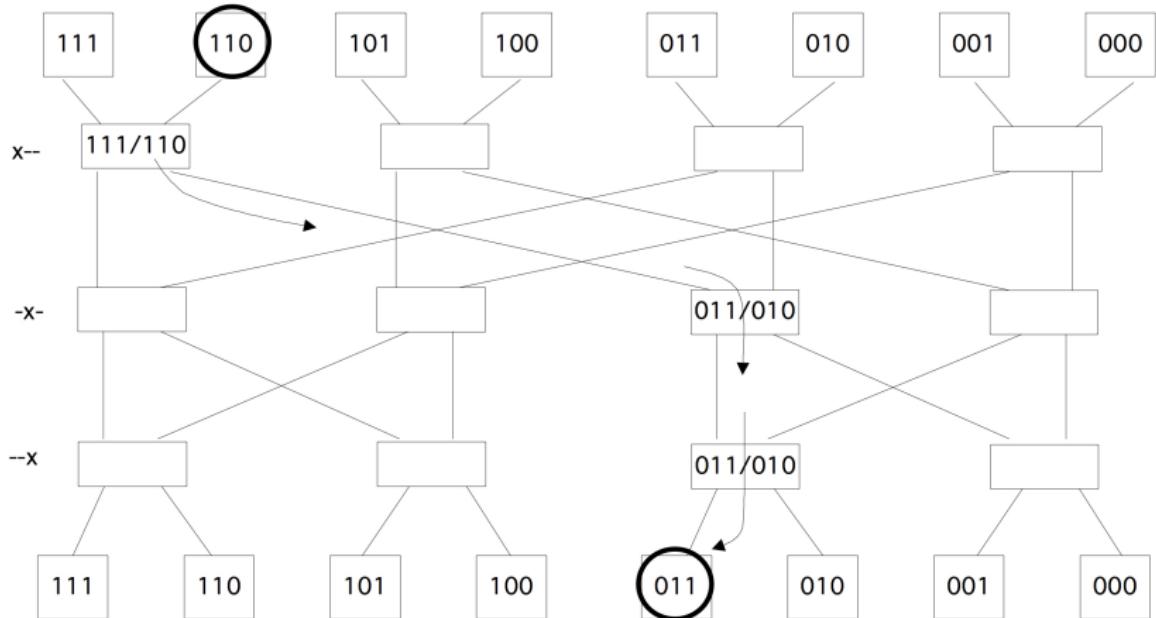


## Uniform memory access

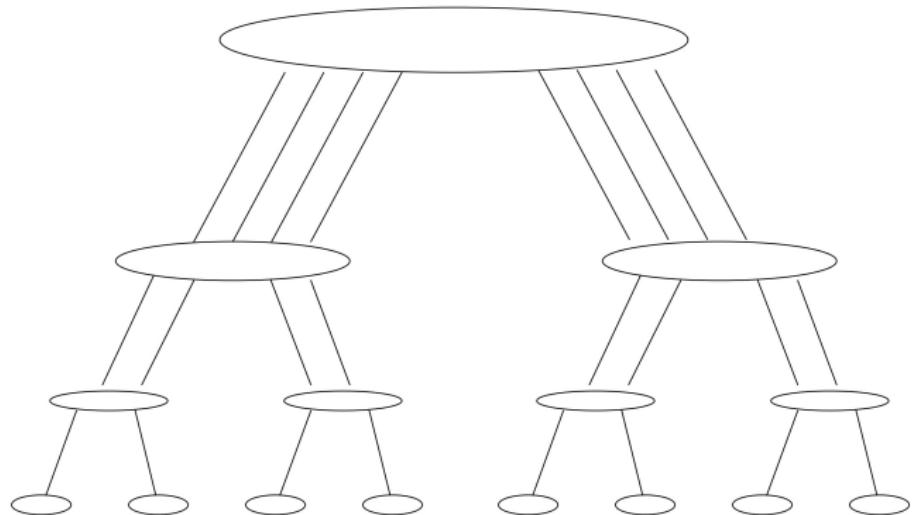


Contention possible

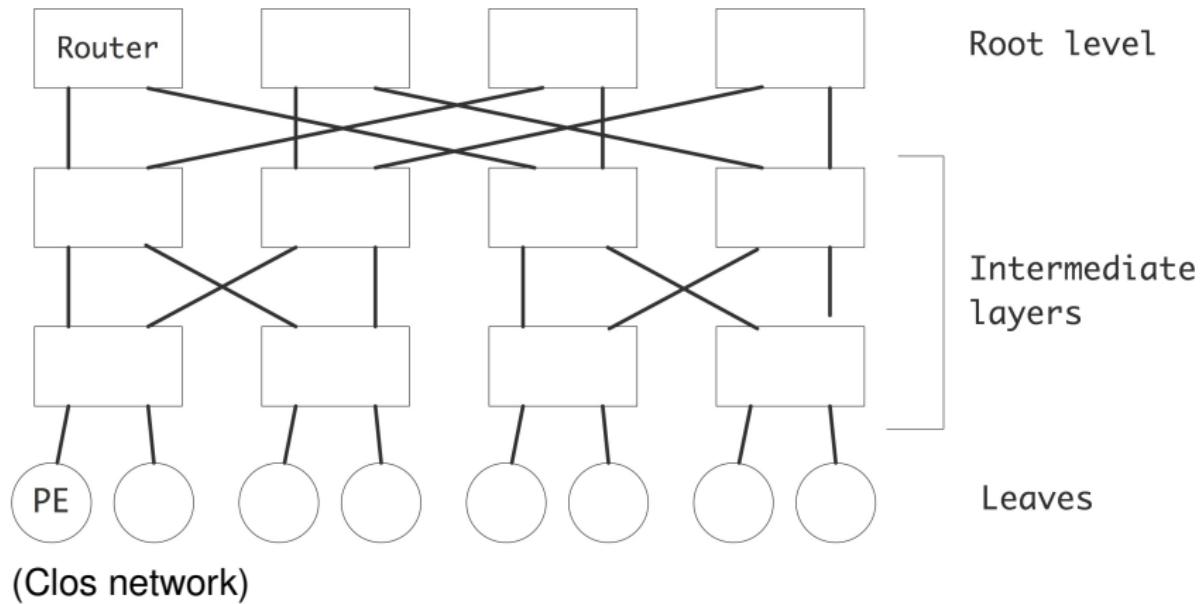
# Route calculation



# Fat Tree



## Fat trees from switching elements



# Fat tree clusters

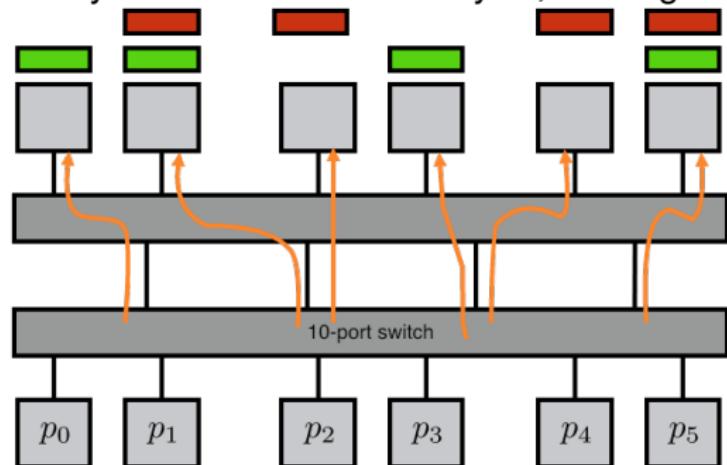


## Exercise 6

Switch contention Suppose the number of processor  $p$  is larger than the number of wires  $w$ .

Write a simulation that investigates the probability of contention if you send  $m \leq w$  message to distinct processors.

Can you do a statistical analysis, starting with a simple case?



## Mesh clusters



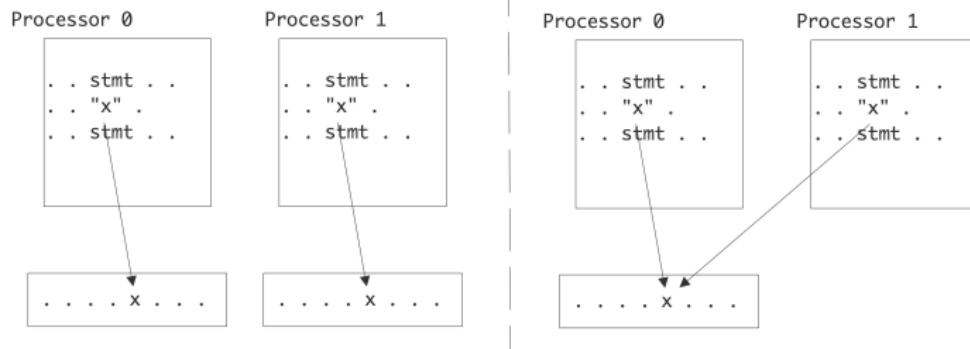
## Levels of locality

- ▶ Core level: private cache, shared cache
- ▶ Node level: numa
- ▶ Network: levels in the switch

# **Programming models**

# Shared vs distributed memory programming

Different memory models:



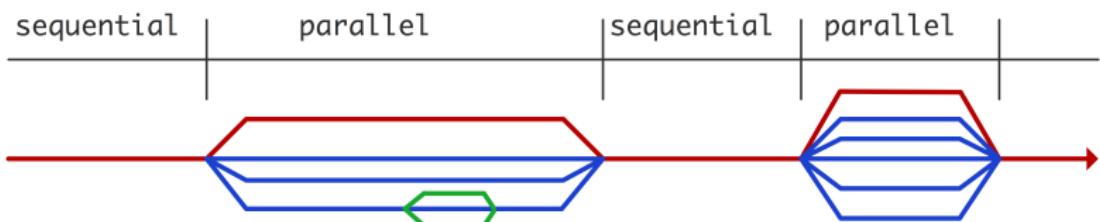
Different questions:

- ▶ Shared memory: synchronization problems such as critical sections
- ▶ Distributed memory: data motion

## **Thread parallelism**

# What is a thread

- ▶ Process: code, heap, stack
- ▶ Thread: same code but private program counter, stack, local variables
- ▶ dynamically (even recursively) created: fork-join



Incremental parallelization!

## Thread context

- ▶ Private data (stack, local variables) is called ‘thread context’
- ▶ Context switch: switch from one thread execution to another
- ▶ context switches are expensive; alternative hyperthreading
- ▶ Intel Xeon Phi: hardware support for 4 threads per core
- ▶ GPUs: fast context switching between many threads

# Thread programming 1

## Pthreads

```
1 pthread_t threads[NTHREADS];
2 printf("forking\n");
3 for (i=0; i<NTHREADS; i++)
4     if (pthread_create(threads+i, NULL, &adder, NULL) !=0)
5         return i+1;
6 printf("joining\n");
7 for (i=0; i<NTHREADS; i++)
8     if (pthread_join(threads[i], NULL) !=0)
9         return NTHREADS+i+1;
```

# Race conditions

Init:  $I=0$

process 1:  $I=I+2$

process 2:  $I=I+3$

scenario 1.	scenario 2.	scenario 3.
	$I = 0$	
read $I = 0$ set $I = 2$ write $I = 2$	read $I = 0$ set $I = 3$ write $I = 3$	read $I = 0$ set $I = 2$ write $I = 2$
		read $I = 2$ set $I = 5$ write $I = 5$
$I = 3$	$I = 2$	$I = 5$

## Dealing with atomic operations

Semaphores, locks, mutexes, critical sections, transactional memory  
Software / hardware

# Cilk

*Sequential code:*

```
int fib(int n){  
    if (n<2) return 1;  
    else {  
        int rst=0;  
        rst += fib(n-1);  
        rst += fib(n-2);  
        return rst;  
    }  
}
```

*Cilk code:*

```
cilk int fib(int n){  
    if (n<2) return 1;  
    else {  
        int rst=0;  
        rst += spawn fib(n-1);  
        rst += spawn fib(n-2);  
        sync;  
        return rst;  
    }  
}
```

Sequential consistency: program output identical to sequential

# OpenMP

- ▶ Directive based
- ▶ Parallel sections, parallel loops, tasks

## **Distributed memory parallelism**

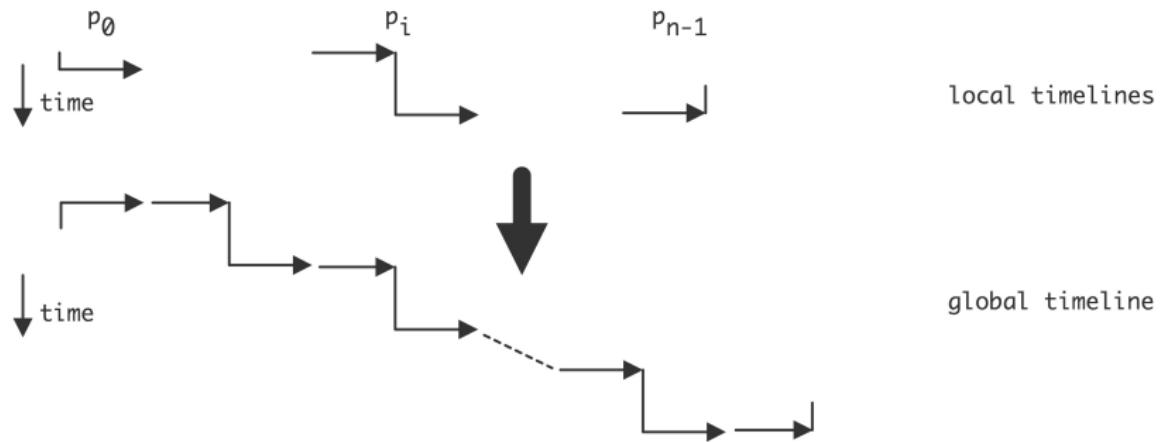
## Global vs local view

$$\begin{cases} y_i \leftarrow y_i + x_{i-1} & i > 0 \\ y_i \text{ unchanged} & i = 0 \end{cases}$$

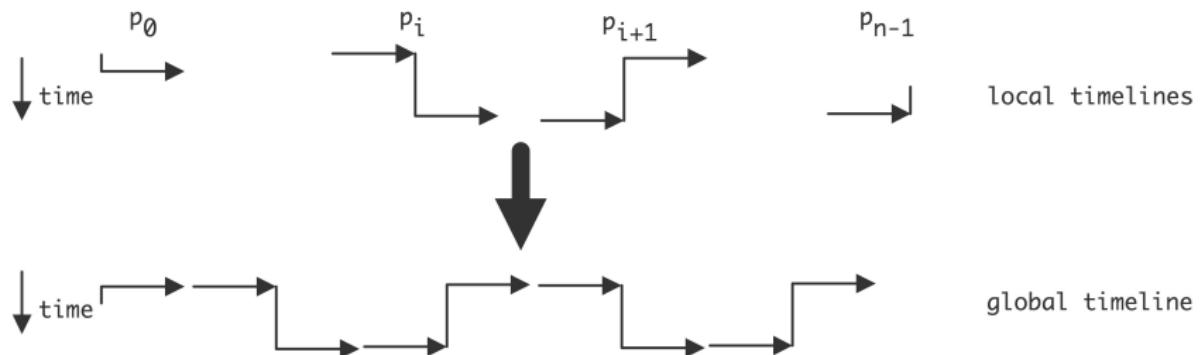
- ▶ If I am processor 0 do nothing, otherwise receive a  $y$  element from the left, add it to my  $x$  element.
- ▶ If I am the last processor do nothing, otherwise send my  $y$  element to the right.

(Let's think this through...)

## Global picture



## Careful coding



## Better approaches

- ▶ Non-blocking send/receive
- ▶ One-sided

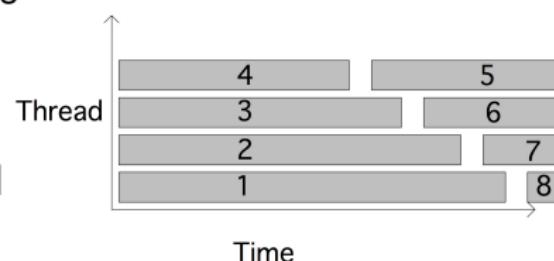
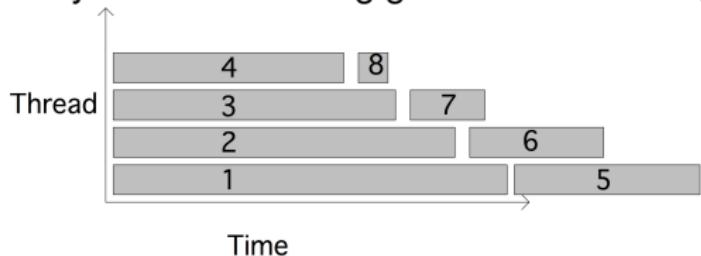
## **Hybrid/heterogeneous parallelism**

## Hybrid computing

- ▶ Use MPI between nodes, OpenMP inside nodes
- ▶ alternative: ignore shared memory and MPI throughout
- ▶ you save: buffers and copying
- ▶ bundling communication, load spread

# Using threads for load balancing

Dynamic scheduling gives load balancing



Hybrid is possible improvement over strict-MPI

## Amdahl's law for hybrid programming

- ▶  $p$  nodes with  $c$  cores each
- ▶  $F_p$  core-parallel fraction, assume full MPI parallel
- ▶ ideal speedup  $pc$ , running time  $T_1/(pc)$ , actually:

$$T_{p,c} = T_1 \left( \frac{F_s}{p} + \frac{F_p}{pc} \right) = \frac{T_1}{pc} (F_s c + F_p) = \frac{T_1}{pc} (1 + F_s(c - 1)).$$

- ▶  $T_1/T_{p,c} \approx p/F_s$
- ▶ Original Amdahl:  $S_p < 1/F_s$ , hybrid programming  $S_p < p/F_s$

## **Design patterns**

# Array of Structures

```
1 struct { int number; double xcoord,ycoord; } _Node;
2 struct { double xtrans,ytrans} _Vector;
3 typedef struct _Node* Node;
4 typedef struct _Vector* Vector;

1 Node *nodes = (node) malloc( n_nodes*sizeof(struct _Node) );
```

# Operations

## Operate

```
1 void shift(node the_point, vector by) {  
2     the_point->xcoord += by->xtrans;  
3     the_point->ycoord += by->ytrans;  
4 }
```

## in a loop

```
1 for (i=0; i<n_nodes; i++) {  
2     shift(nodes[i],shift_vector);  
3 }
```

# Along come the 80s

## Vector operations

```
1 node_numbers = (int*) malloc( n_nodes*sizeof(int) );
2 node_xcoords = // et cetera
3 node_ycoords = // et cetera
```

and you would iterate

```
1 for (i=0; i<n_nodes; i++) {
2     node_xoords[i] += shift_vector->xtrans;
3     node_yoords[i] += shift_vector->ytrans;
4 }
```

and the wheel of reinvention turns further

The original design was better for MPI in the 1990s  
except when vector instructions (and GPUs) came along in the 2000s

## Latency hiding

- ▶ Memory and network are slow, prevent having to wait for it
- ▶ Hardware magic: out-of-order execution, caches, prefetching

## Explicit latency hiding

Matrix vector product

$$\forall_{i \in I_p} : y_i = \sum_j a_{ij} x_j.$$

$x$  needs to be gathered:

$$\forall_{i \in I_p} : y_i = \left( \sum_{j \text{ local}} + \sum_{j \text{ not local}} \right) a_{ij} x_j.$$

Overlap loads and local operations

Possible in MPI and Xeon Phi offloading,  
very hard to do with caches

**What's left**

## Parallel languages

- ▶ Co-array Fortran: extensions to the Fortran standard
- ▶ X10
- ▶ Chapel
- ▶ UPC
- ▶ BSP
- ▶ MapReduce
- ▶ Pregel, ...

## UPC example

```
1 #define N 100*THREADS
2
3 shared int v1[N], v2[N], v1plusv2[N];
4
5 void main()
6 {
7     int i;
8     upc_forall(i=0; i<N; i++)
9         v1plusv2[i]=v1[i]+v2[i];
10 }
```

## Co-array Fortran example

Explicit dimension for ‘images’:

```
1 Real, dimension(100), codimension[*] :: X
2 Real :: X(100)[*]
3 Real :: X(100,200)[10,0:9,*]
```

determined by runtime environment

## Grab bag of other approaches

- ▶ OS-based: data movement induced by cache misses
- ▶ Active messages: application level Remote Procedure Call  
(see: Charm++)

**Load balancing, locality, space-filling curves**

## The load balancing problem

- ▶ Application load can change dynamically
  - e.g., mesh refinement, time-dependent problems
- ▶ Splitting off and merging loads
- ▶ No real software support: write application anticipating load management
- ▶ Initial balancing: graph partitioners

## Load balancing and performance

- ▶ Assignment to arbitrary processor violates locality
- ▶ Need a dynamic load assignment scheme that preserves locality under load migration
- ▶ Fairly easy for regular problems, for irregular?

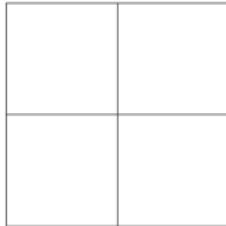
## **Space-filling curves**

# Adaptive refinement and load assignment

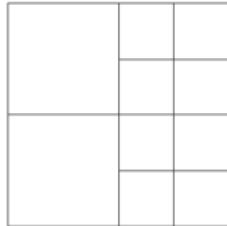
Level 0



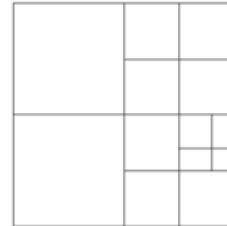
Level 1



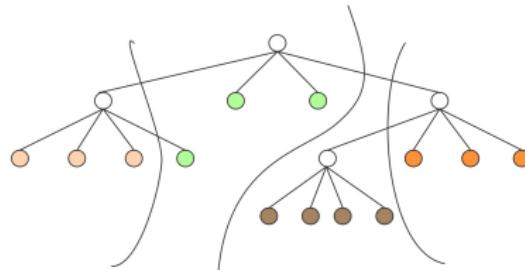
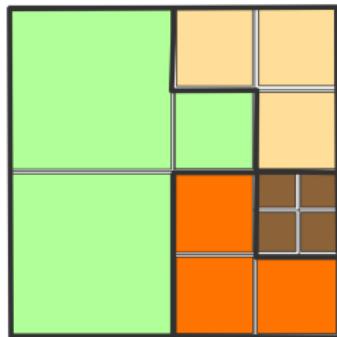
Level 2



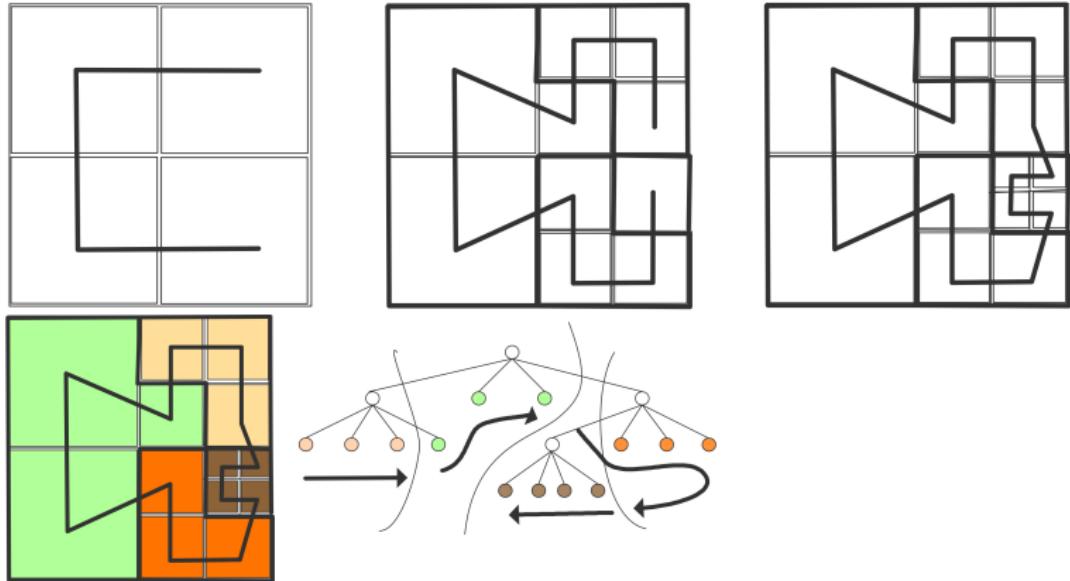
Level 3



...



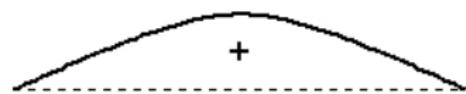
# Assignment through Space-Filling Curve



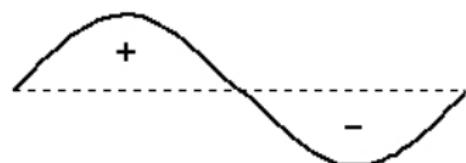
## **Domain partitioning by Fiedler vectors**

# Inspiration from physics

## Modes of a Vibrating String



Lowest Frequency  $\lambda(1)$



Second Frequency  $\lambda(2)$



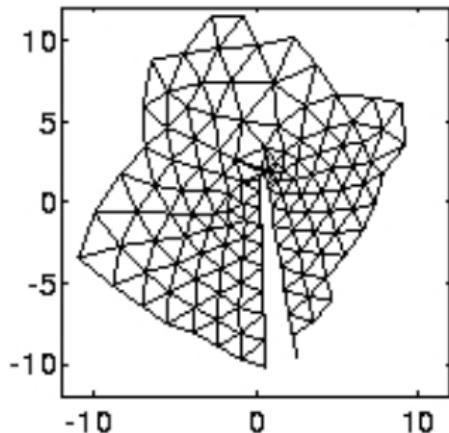
Third Frequency  $\lambda(3)$

## Graph laplacian

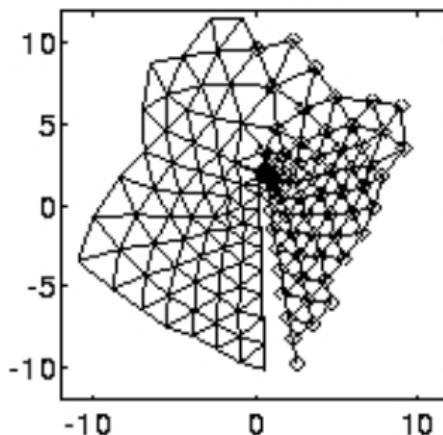
- ▶ Set  $G_{ij} = -1$  if edge  $(i,j)$
- ▶ Set  $G_{ii}$  positive to give zero rowsums
- ▶ First eigenvector is zero, positive eigenvector
- ▶ Second eigenvector has pos/neg, divides in two
- ▶  $n$ -th eigenvector divides in  $n$  parts

## Fiedler in a picture

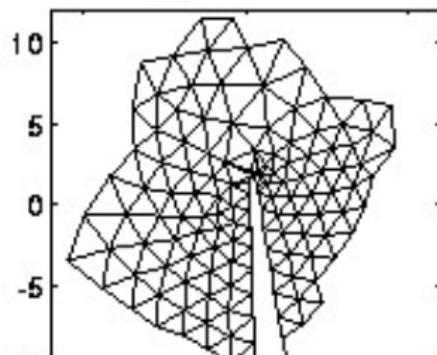
Original FE mesh



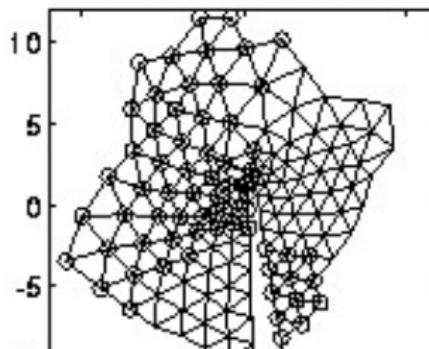
Circle node  $i$  if  $v_2(i) > 0$



Original FE mesh



Circle node  $i$  if  $v_4(i) > 0$



# **Collectives**

**Collectives as building blocks; complexity**

## Collectives

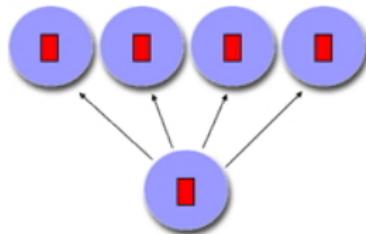
Gathering and spreading information:

- ▶ Every process has data, you want to bring it together;
- ▶ One process has data, you want to spread it around.

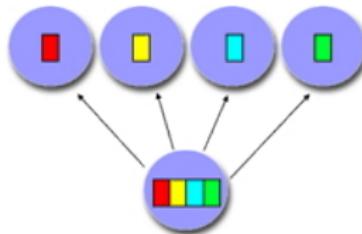
Root process: the one doing the collecting or disseminating.

Basic cases:

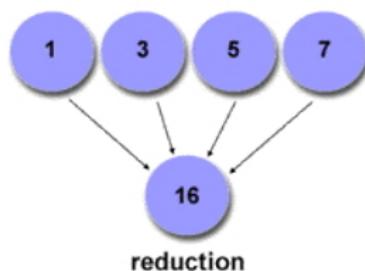
- ▶ Collect data: gather.
- ▶ Collect data and compute some overall value (sum, max): reduction.
- ▶ Send the same data to everyone: broadcast.
- ▶ Send individual data to each process: scatter.



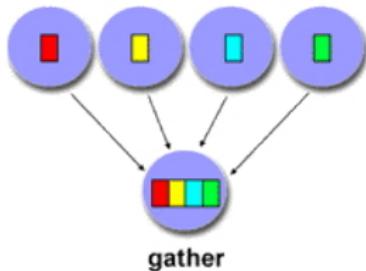
broadcast



scatter



reduction



gather

## Collective scenarios

How would you realize the following scenarios with collectives?

- ▶ Let each process compute a random number. You want to print the maximum of these numbers to your screen.
- ▶ Each process computes a random number again. Now you want to scale these numbers by their maximum.
- ▶ Let each process compute a random number. You want to print on what processor the maximum value is computed.

## Simple model of parallel computation

- ▶  $\alpha$ : message latency
- ▶  $\beta$ : time per word (inverse of bandwidth)
- ▶  $\gamma$ : time per floating point operation

Send  $n$  items and do  $m$  operations:

$$cost = \alpha + \beta \cdot n + \gamma \cdot m$$

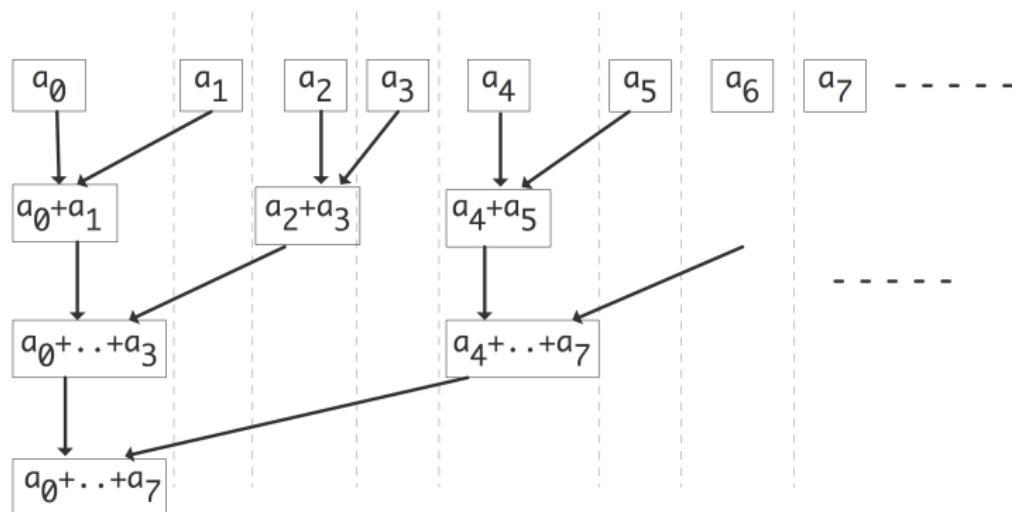
Pure sends: no  $\gamma$  term,

pure computation: no  $\alpha, \beta$  terms,

sometimes mixed: reduction

## Model for collectives

- ▶ One simultaneous send and receive:
- ▶ doubling of active processors
- ▶ collectives have a  $\alpha \log_2 p$  cost component



## Broadcast

	$t = 0$	$t = 1$	$t = 2$
$p_0$	$x_0 \downarrow, x_1 \downarrow, x_2 \downarrow, x_3 \downarrow$	$x_0 \downarrow, x_1 \downarrow, x_2 \downarrow, x_3 \downarrow$	$x_0, x_1, x_2, x_3$
$p_1$		$x_0 \downarrow, x_1 \downarrow, x_2 \downarrow, x_3 \downarrow$	$x_0, x_1, x_2, x_3$
$p_2$			$x_0, x_1, x_2, x_3$
$p_3$			$x_0, x_1, x_2, x_3$

On  $t = 0$ ,  $p_0$  sends to  $p_1$ ; on  $t = 1$   $p_0, p_1$  send to  $p_2, p_3$ .

Optimal complexity:

$$\lceil \log_2 p \rceil \alpha + n\beta.$$

Actual complexity:

$$\lceil \log_2 p \rceil (\alpha + n\beta).$$

Good enough for short vectors.

## Long vector broadcast

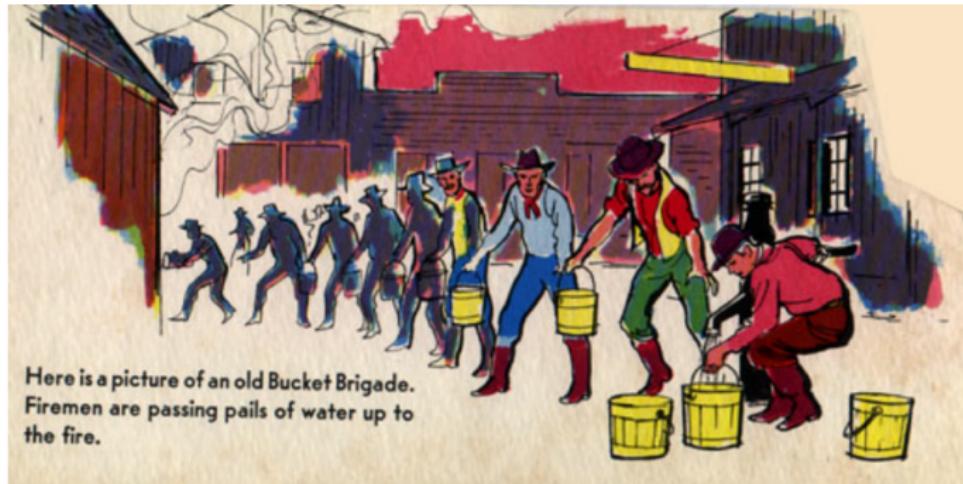
Start with a scatter:

	$t = 0$	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0 \downarrow, x_1, x_2, x_3$	$x_0, x_1 \downarrow, x_2, x_3$	$x_0, x_1, x_2 \downarrow, x_3$	$x_0, x_1, x_2, x_3 \downarrow$
$p_1$		$x_1$		
$p_2$			$x_2$	
$p_3$				$x_3$

takes  $p - 1$  messages of size  $N/p$ , for a total time of

$$T_{\text{scatter}}(N, P) = (p - 1)\alpha + (p - 1) \cdot \frac{N}{p} \cdot \beta.$$

# Bucket brigade



Here is a picture of an old Bucket Brigade.  
Firemen are passing pails of water up to  
the fire.

## Long vector broadcast

After the scatter do a bucket-allgather:

	$t = 0$	$t = 1$	<i>etcetera</i>
$p_0$	$x_0 \downarrow$	$x_0$	$x_3 \downarrow$ $x_0, x_2, x_3$
$p_1$	$x_1 \downarrow$	$x_0 \downarrow, x_1$	$x_0, x_1, x_3$
$p_2$	$x_2 \downarrow$	$x_1 \downarrow, x_2$	$x_0, x_1, x_2$
$p_3$	$x_3 \downarrow$	$x_2 \downarrow, x_3$	$x_1, x_2, x_3$

Each partial message gets sent  $p - 1$  times, so this stage also has a complexity of

$$T_{\text{bucket}}(N, P) = (p - 1)\alpha + (p - 1) \cdot \frac{N}{p} \cdot \beta.$$

Better if  $N$  large.

# Reduce

Optimal complexity:

$$\lceil \log_2 p \rceil \alpha + n\beta + \frac{p-1}{p} \gamma n.$$

Spanning tree algorithm:

	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0^{(0)}, x_1^{(0)}, x_2^{(0)}, x_3^{(0)}$	$x_0^{(0:1)}, x_1^{(0:1)}, x_2^{(0:1)}, x_3^{(0:1)}$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
$p_1$	$x_0^{(1)} \uparrow, x_1^{(1)} \uparrow, x_2^{(1)} \uparrow, x_3^{(1)} \uparrow$		
$p_2$	$x_0^{(2)}, x_1^{(2)}, x_2^{(2)}, x_3^{(2)}$	$x_0^{(2:3)} \uparrow, x_1^{(2:3)} \uparrow, x_2^{(2:3)} \uparrow, x_3^{(2:3)} \uparrow$	
$p_3$	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)} \uparrow, x_3^{(3)} \uparrow$		

Running time

$$\lceil \log_2 p \rceil (\alpha + n\beta + \frac{p-1}{p} \gamma n).$$

Good enough for short vectors.

# Allreduce

Allreduce  $\equiv$  Reduce + Broadcast

	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0^{(0)} \downarrow, x_1^{(0)} \downarrow, x_2^{(0)} \downarrow, x_3^{(0)} \downarrow$	$x_0^{(0:1)} \downarrow\downarrow, x_1^{(0:1)} \downarrow\downarrow, x_2^{(0:1)} \downarrow\downarrow, x_3^{(0:1)} \downarrow\downarrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
$p_1$	$x_0^{(1)} \uparrow, x_1^{(1)} \uparrow, x_2^{(1)} \uparrow, x_3^{(1)} \uparrow$	$x_0^{(0:1)} \downarrow\downarrow, x_1^{(0:1)} \downarrow\downarrow, x_2^{(0:1)} \downarrow\downarrow, x_3^{(0:1)} \downarrow\downarrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
$p_2$	$x_0^{(2)} \downarrow, x_1^{(2)} \downarrow, x_2^{(2)} \downarrow, x_3^{(2)} \downarrow$	$x_0^{(2:3)} \uparrow\uparrow, x_1^{(2:3)} \uparrow\uparrow, x_2^{(2:3)} \uparrow\uparrow, x_3^{(2:3)} \uparrow\uparrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$
$p_3$	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)} \uparrow, x_3^{(3)} \uparrow$	$x_0^{(2:3)} \uparrow\uparrow, x_1^{(2:3)} \uparrow\uparrow, x_2^{(2:3)} \uparrow\uparrow, x_3^{(2:3)} \uparrow\uparrow$	$x_0^{(0:3)}, x_1^{(0:3)}, x_2^{(0:3)}, x_3^{(0:3)}$

Same running time as regular reduce!

## Allgather

Gather  $n$  elements: each processor owns  $n/p$ ;  
optimal running time

$$\lceil \log_2 p \rceil \alpha + \frac{p-1}{p} n \beta.$$

	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0 \downarrow$	$x_0x_1 \downarrow$	$x_0x_1x_2x_3$
$p_1$	$x_1 \uparrow$	$x_0x_1 \downarrow$	$x_0x_1x_2x_3$
$p_2$	$x_2 \downarrow$	$x_2x_3 \uparrow$	$x_0x_1x_2x_3$
$p_3$	$x_3 \uparrow$	$x_2x_3 \uparrow$	$x_0x_1x_2x_3$

Same time as gather, half of gather-and-broadcast.

## Reduce-scatter

	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0^{(0)}, x_1^{(0)}, x_2^{(0)} \downarrow, x_3^{(0)} \downarrow$	$x_0^{(0:2:2)}, x_1^{(0:2:2)} \downarrow$	$x_0^{(0:3)}$
$p_1$	$x_0^{(1)}, x_1^{(1)}, x_2^{(1)} \downarrow, x_3^{(1)} \downarrow$	$x_0^{(1:3:2)} \uparrow, x_1^{(1:3:2)}$	$x_1^{(0:3)}$
$p_2$	$x_0^{(2)} \uparrow, x_1^{(2)} \uparrow, x_2^{(2)}, x_3^{(2)}$	$x_2^{(0:2:2)}, x_3^{(0:2:2)} \downarrow$	$x_2^{(0:3)}$
$p_3$	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)}, x_3^{(3)}$	$x_0^{(1:3:2)} \uparrow, x_1^{(1:3:2)}$	$x_3^{(0:3)}$

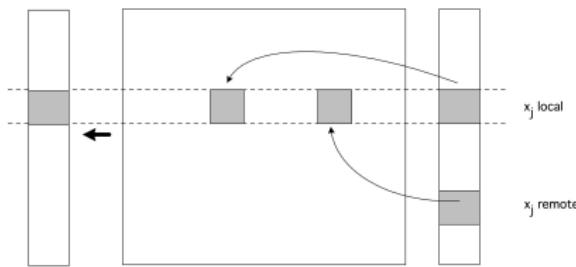
$$\lceil \log_2 p \rceil \alpha + \frac{p-1}{p} n(\beta + \gamma).$$

# **Applications**

# **Scalability analysis of dense matrix-vector product**

## Parallel matrix-vector product; general

- ▶ Assume a division by block rows
- ▶ Every processor  $p$  has a set of row indices  $I_p$

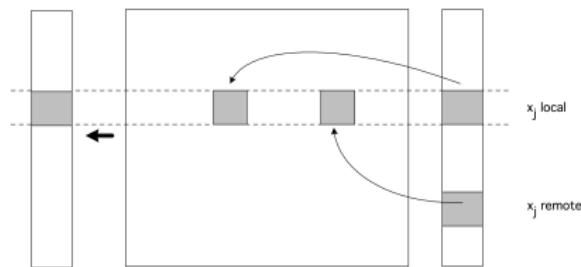


Mvp on processor  $p$ :

$$\forall i \in I_p : y_i = \sum_j a_{ij} x_j = \sum_q \sum_{j \in I_q} a_{ij} x_j$$

# Local and remote operations

Local and remote parts:



$$\forall_{i \in I_p} : y_i = \sum_{j \in I_p} a_{ij} x_j + \sum_{q \neq p} \sum_{j \in I_q} a_{ij} x_j$$

Local part  $I_p$  can be executed right away,  $I_q$  requires communication.

## How to deal with remote parts

- ▶ Very flexible: mix of working on local parts, and receiving remote parts.
- ▶ More orchestrated:
  1. each process gets a full copy of the input vector (how?)
  2. then operates on the whole input

Compare?

(Are we making a big assumption here?)

## Dense MVP

- ▶ Separate communication and computation:
- ▶ first allgather
- ▶ then matrix-vector product

## Cost computation 1.

Algorithm:

---

Step	Cost (lower bound)
Allgather $x_i$ so that $x$ is available on all nodes	
Locally compute $y_i = A_i x$	$\approx 2 \frac{n^2}{P} \gamma$

---

## Allgather

Assume that data arrives over a binary tree:

- ▶ latency  $\alpha \log_2 P$
- ▶ transmission time, receiving  $n/P$  elements from  $P - 1$  processors

Algorithm with cost:

---

Step	Cost (lower bound)
Allgather $x_i$ so that $x$ is available on all nodes	$\lceil \log_2(P) \rceil \alpha + \frac{P-1}{P} n\beta \approx \log_2(P)\alpha + n\beta$
Locally compute $y_i = A_i x$	$\approx 2\frac{n^2}{P}\gamma$

---

## Parallel efficiency

Speedup:

$$\begin{aligned} S_p^{\text{1D-row}}(n) &= \frac{T_1(n)}{T_p^{\text{1D-row}}(n)} \\ &= \frac{2n^2\gamma}{2\frac{n^2}{p}\gamma + \log_2(p)\alpha + n\beta} \\ &= \frac{p}{1 + \frac{p\log_2(p)}{2n^2}\frac{\alpha}{\gamma} + \frac{p}{2n}\frac{\beta}{\gamma}} \end{aligned}$$

Efficiency:

$$\begin{aligned} E_p^{\text{1D-row}}(n) &= \frac{S_p^{\text{1D-row}}(n)}{p} \\ &= \frac{1}{1 + \frac{p\log_2(p)}{2n^2}\frac{\alpha}{\gamma} + \frac{p}{2n}\frac{\beta}{\gamma}}. \end{aligned}$$

Strong scaling, weak scaling?

## Optimistic scaling

Processors fixed, problem grows:

$$E_p^{\text{1D-row}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}.$$

Roughly  $E_p \sim 1 - n^{-1}$

## Strong scaling

Problem fixed,  $p \rightarrow \infty$

$$E_p^{\text{1D-row}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}.$$

## Strong scaling

Problem fixed,  $p \rightarrow \infty$

$$E_p^{\text{1D-row}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}.$$

Roughly  $E_p \sim p^{-1}$

## Weak scaling

Memory fixed:

$$M = n^2/p$$

$$E_p^{1\text{D-row}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}} = \frac{1}{1 + \frac{\log_2(p)}{2M} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2\sqrt{M}} \frac{\beta}{\gamma}}$$

## Weak scaling

Memory fixed:

$$M = n^2/p$$

$$E_p^{1\text{D-row}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}} = \frac{1}{1 + \frac{\log_2(p)}{2M} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2\sqrt{M}} \frac{\beta}{\gamma}}$$

Does not scale:  $E_p \sim 1/\sqrt{p}$

problem in  $\beta$  term: too much communication

## Two-dimensional partitioning

$x_0$	$a_{00}$	$a_{01}$	$a_{02}$	$y_0$	$a_{03}$	$a_{04}$	$a_{05}$	$y_1$	$a_{06}$	$a_{07}$	$a_{08}$	$y_2$	$a_{09}$	$a_{0,10}$	$a_{0,11}$
$a_{10}$	$a_{11}$	$a_{12}$			$a_{13}$	$a_{14}$	$a_{15}$		$a_{16}$	$a_{17}$	$a_{18}$		$a_{19}$	$a_{1,10}$	$a_{1,11}$
$a_{20}$	$a_{21}$	$a_{22}$			$a_{23}$	$a_{24}$	$a_{25}$		$a_{26}$	$a_{27}$	$a_{28}$		$a_{29}$	$a_{2,10}$	$a_{2,11}$
$a_{30}$	$a_{31}$	$a_{32}$			$a_{33}$	$a_{34}$	$a_{35}$		$a_{37}$	$a_{37}$	$a_{38}$		$a_{39}$	$a_{3,10}$	$a_{3,11}$
$x_1$					$x_4$				$x_7$				$x_{10}$		
$a_{40}$	$a_{41}$	$a_{42}$	$y_4$		$a_{43}$	$a_{44}$	$a_{45}$		$a_{46}$	$a_{47}$	$a_{48}$		$a_{49}$	$a_{4,10}$	$a_{4,11}$
$a_{50}$	$a_{51}$	$a_{52}$			$a_{53}$	$a_{54}$	$a_{55}$	$y_5$	$a_{56}$	$a_{57}$	$a_{58}$		$a_{59}$	$a_{5,10}$	$a_{5,11}$
$a_{60}$	$a_{61}$	$a_{62}$			$a_{63}$	$a_{64}$	$a_{65}$		$a_{66}$	$a_{67}$	$a_{68}$	$y_6$	$a_{69}$	$a_{6,10}$	$a_{6,11}$
$a_{70}$	$a_{71}$	$a_{72}$			$a_{73}$	$a_{74}$	$a_{75}$		$a_{77}$	$a_{77}$	$a_{78}$		$a_{79}$	$a_{7,10}$	$a_{7,11}$
$x_2$					$x_5$				$x_8$				$x_{11}$		
$a_{80}$	$a_{81}$	$a_{82}$	$y_8$		$a_{83}$	$a_{84}$	$a_{85}$		$a_{86}$	$a_{87}$	$a_{88}$		$a_{89}$	$a_{8,10}$	$a_{8,11}$
$a_{90}$	$a_{91}$	$a_{92}$			$a_{93}$	$a_{94}$	$a_{95}$	$y_9$	$a_{96}$	$a_{97}$	$a_{98}$		$a_{99}$	$a_{9,10}$	$a_{9,11}$
$a_{10,0}$	$a_{10,1}$	$a_{10,2}$			$a_{10,3}$	$a_{10,4}$	$a_{10,5}$		$a_{10,6}$	$a_{10,7}$	$a_{10,8}$	$y_{10}$	$a_{10,9}$	$a_{10,10}$	$a_{10,11}$
$a_{11,0}$	$a_{11,1}$	$a_{11,2}$			$a_{11,3}$	$a_{11,4}$	$a_{11,5}$		$a_{11,7}$	$a_{11,7}$	$a_{11,8}$		$a_{11,9}$	$a_{11,10}$	$a_{11,11}$

## Two-dimensional partitioning

Processor grid  $p = r \times c$ , assume  $r, c \approx \sqrt{p}$ .

$x_0$ $a_{00}$ $a_{10}$ $a_{20}$ $a_{30}$	$a_{01}$ $a_{11}$ $a_{21}$ $a_{31}$	$a_{02}$ $a_{12}$ $a_{22}$ $a_{32}$	$y_0$	$x_3$  $y_1$	$x_6$  $y_2$	$x_9$  $y_3$
$x_1 \uparrow$  $y_4$		$x_4$  $y_5$		$x_7$  $y_6$		$x_{10}$  $y_7$
$x_2 \uparrow$  $y_8$		$x_5$  $y_9$		$x_8$  $y_{10}$		$x_{11}$  $y_{11}$

## Key to the algorithm

- ▶ Consider block  $(i, j)$
- ▶ it needs to multiply by the  $xs$  in column  $j$
- ▶ it produces part of the result of row  $i$

## Algorithm

- ▶ Collecting  $x_j$  on each processor  $p_{ij}$  by an *allgather* inside the processor columns.
- ▶ Each processor  $p_{ij}$  then computes  $y_{ij} = A_{ij}x_j$ .
- ▶ Gathering together the pieces  $y_{ij}$  in each processor row to form  $y_i$ , distribute this over the processor row: combine to form a *reduce-scatter*.
- ▶ Setup for the next  $A$  or  $A^t$  product

## Analysis 1.

---

Step	Cost (lower bound)
Allgather $x_i$ 's within columns	$\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n\beta$ $\approx \log_2(r)\alpha + \frac{n}{c}\beta$
Perform local matrix-vector multiply	$\approx 2\frac{n^2}{p}\gamma$
Reduce-scatter $y_i$ 's within rows	

---

## Reduce-scatter

Time:

$$\lceil \log_2 p \rceil \alpha + \frac{p-1}{p} n(\beta + \gamma).$$

Step	Cost (lower bound)
Allgather $x_i$ 's within columns	$\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n\beta$ $\approx \log_2(r)\alpha + \frac{n}{c}\beta$
Perform local matrix-vector multiply	$\approx 2\frac{n^2}{p}\gamma$
Reduce-scatter $y_i$ 's within rows	$\lceil \log_2(c) \rceil \alpha + \frac{c-1}{p} n\beta + \frac{c-1}{p} m\gamma$ $\approx \log_2(c)\alpha + \frac{n}{r}\beta + \frac{n}{r}\gamma$

## Efficiency

Let  $r = c = \sqrt{p}$ , then

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta + \gamma)}{\gamma}}$$

## Strong scaling

Same story as before for  $p \rightarrow \infty$ :

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta + \gamma)}{\gamma}} \sim p^{-1}$$

No strong scaling

## Weak scaling

Constant memory  $M = n^2/p$ :

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta + \gamma)}{\gamma}}$$

## Weak scaling

Constant memory  $M = n^2/p$ :

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta+\gamma)}{\gamma}} = \frac{1}{1 + \frac{\log_2(p)}{2M} \frac{\alpha}{\gamma} + \frac{1}{2\sqrt{M}} \frac{(2\beta+\gamma)}{\gamma}}$$

## Weak scaling

Constant memory  $M = n^2/p$ :

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta+\gamma)}{\gamma}} = \frac{1}{1 + \frac{\log_2(p)}{2M} \frac{\alpha}{\gamma} + \frac{1}{2\sqrt{M}} \frac{(2\beta+\gamma)}{\gamma}}$$

Weak scaling:

for  $p \rightarrow \infty$  this is  $\approx 1/\log_2 p$ :

only slowly decreasing.

## LU factorizations

- ▶ Needs a cyclic distribution
- ▶ This is very hard to program, so:
- ▶ Scalapack, 1990s product, not extendible, impossible interface
- ▶ Elemental: 2010s product, extendible, nice user interface (and it is way faster)

# **Sparse linear algebra**

## Boundary value problems

Consider in 1D

$$\begin{cases} -u''(x) = f(x, u, u') & x \in [a, b] \\ u(a) = u_a, u(b) = u_b \end{cases}$$

in 2D:

$$\begin{cases} -u_{xx}(\bar{x}) - u_{yy}(\bar{x}) = f(\bar{x}) & \bar{x} \in \Omega = [0, 1]^2 \\ u(\bar{x}) = u_0 & \bar{x} \in \delta\Omega \end{cases}$$

## Approximation of 2nd order derivatives

Taylor series (write  $h$  for  $\delta x$ ):

$$u(x+h) = u(x) + u'(x)h + u''(x)\frac{h^2}{2!} + u'''(x)\frac{h^3}{3!} + u^{(4)}(x)\frac{h^4}{4!} + u^{(5)}(x)\frac{h^5}{5!} + \dots$$

and

$$u(x-h) = u(x) - u'(x)h + u''(x)\frac{h^2}{2!} - u'''(x)\frac{h^3}{3!} + u^{(4)}(x)\frac{h^4}{4!} - u^{(5)}(x)\frac{h^5}{5!} + \dots$$

Subtract:

$$u(x+h) + u(x-h) = 2u(x) + u''(x)h^2 + u^{(4)}(x)\frac{h^4}{12} + \dots$$

so

$$u''(x) = \frac{u(x+h) - 2u(x) + u(x-h)}{h^2} - u^{(4)}(x)\frac{h^4}{12} + \dots$$

Numerical scheme:

$$-\frac{u(x+h) - 2u(x) + u(x-h)}{h^2} = f(x, u(x), u'(x))$$

(2nd order PDEs are very common!)

This leads to linear algebra

$$-u_{xx} = f \rightarrow \frac{2u(x) - u(x+h) - u(x-h)}{h^2} = f(x, u(x), u'(x))$$

Equally spaced points on  $[0, 1]$ :  $x_k = kh$  where  $h = 1/(n+1)$ , then

$$-u_{k+1} + 2u_k - u_{k-1} = -1/h^2 f(x_k, u_k, u'_k) \quad \text{for } k = 1, \dots, n$$

Written as matrix equation:

$$\begin{pmatrix} 2 & -1 & & \\ -1 & 2 & -1 & \\ & \ddots & \ddots & \ddots \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ \vdots \end{pmatrix} = \begin{pmatrix} f_1 + u_0 \\ f_2 \\ \vdots \end{pmatrix}$$

## Matrix properties

- ▶ Very sparse, banded
- ▶ Symmetric (only because 2nd order problem)
- ▶ Sign pattern: positive diagonal, nonpositive off-diagonal  
(true for many second order methods)
- ▶ Positive definite (just like the continuous problem)
- ▶ Constant diagonals (from constant coefficients in the DE)

## Sparse matrix in 2D case

Sparse matrices so far were tridiagonal: only in 1D case.

Two-dimensional:  $-u_{xx} - u_{yy} = f$  on unit square  $[0, 1]^2$

Difference equation:

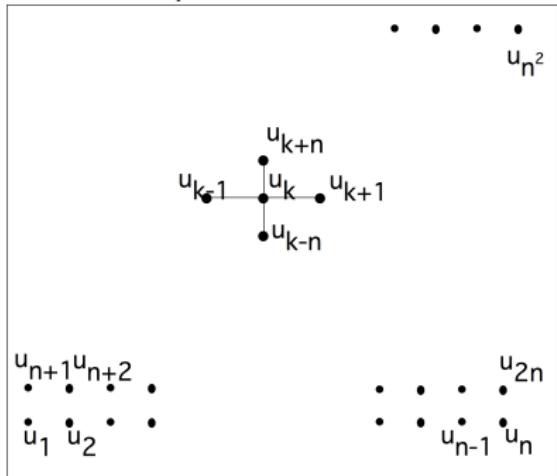
$$4u(x, y) - u(x + h, y) - u(x - h, y) - u(x, y + h) - u(x, y - h) = h^2 f(x, y)$$

$$4u_k - u_{k-1} - u_{k+1} - u_{k-n} - u_{k+n} = f_k$$

Consider a graph where  $\{u_k\}_k$  are the edges  
and  $(u_i, u_j)$  is an edge iff  $a_{ij} \neq 0$ .

# The graph view of things

Poisson eq:



This is a graph!

This is the (adjacency) graph of a sparse matrix.

## Sparse matrix from 2D equation

$$\left( \begin{array}{cccc|ccc|c} 4 & -1 & & 0 & -1 & & & 0 \\ -1 & 4 & 1 & & -1 & & & \\ \ddots & \ddots & \ddots & & \ddots & & & \\ & \ddots & \ddots & -1 & & & & \\ 0 & & -1 & 4 & 0 & & & -1 \\ \hline -1 & & & 0 & 4 & -1 & & -1 \\ & -1 & & & -1 & 4 & -1 & -1 \\ & \uparrow & \ddots & & \uparrow & \uparrow & \uparrow & \uparrow \\ k-n & & & & k-1 & k & k+1 & k+n \\ \hline & & & -1 & & & -1 & \\ & & & & \ddots & & & \ddots \\ \end{array} \right)$$

## Matrix properties

- ▶ Very sparse, banded
- ▶ Factorization takes less than  $n^2$  space,  $n^3$  work
- ▶ Symmetric (only because 2nd order problem)
- ▶ Sign pattern: positive diagonal, nonpositive off-diagonal  
(true for many second order methods)
- ▶ Positive definite (just like the continuous problem)
- ▶ Constant diagonals: only because of the constant coefficient differential equation
- ▶ Factorization: lower complexity than dense, recursion length less than  $N$ .

## Realistic meshes

